

# The Due Process Deficit: Auditing AI Governance in U.S. Higher Education

MARIA TELEKI, Texas A&M University, U.S.A.

ANNA SEO GYEONG CHOI\*, Cornell University, U.S.A.

ANNE DURAY\*, AAAS Science and Technology Policy Fellow, U.S.A.

HAORAN LIU, Texas A&M University, U.S.A.

JUNYAN ZHANG, Independent, U.S.A.

XIANGJUE DONG, Texas A&M University, U.S.A.

DILMA DA SILVA, Texas A&M University, U.S.A.

ALLISON KOENECKE<sup>†</sup>, Cornell Tech, U.S.A.

JAMES CAVERLEE<sup>†</sup>, Texas A&M University, U.S.A.

U.S. universities have met Artificial Intelligence (AI) with policy – revised honor codes, AI use guidelines, and detection tools to enforce them – but the procedural infrastructure that would let affected parties understand, contest, and seek redress for AI-related decisions has not kept pace. We present ACAI-US79, a benchmark dataset allowing for institutional auditing of AI governance across 79 U.S. universities spanning R1, R2, and teaching-focused institutions, and ACAI, the Academic AI Capacity Index, an interpretable measure of the public legibility of governance structures that allocate authority and accountability. Unlike prior audits focused on pedagogical guidance at elite or R1-only samples, ACAI shifts the analytic frame to procedural accountability, evaluating four domains – policy clarity, faculty support, feedback mechanisms, and AI detection tool governance – through time-bounded review of institutionally authoritative materials. We find that governance capacity is systematically concentrated in rule articulation while mechanisms for participation, feedback, and procedural constraint are underdeveloped, and that AI research intensity is weakly correlated with AI governance capacity – i.e., we find that technical leadership does not translate into procedural accountability. We further show that LLM-based audits produce unstable institutional rankings, demonstrating the continued necessity of human judgment in high-stakes audit contexts. We release the dataset, audit instrument, and public website at <http://acai-us79.org/> to support transparency, critique, and institutional self-reflection, laying empirical groundwork for accountable AI governance in higher education.

CCS Concepts: • **Applied computing** → **Education**; **Law**; *Computer-managed instruction*.

Additional Key Words and Phrases: AI, university, policy, governance

\*Both authors contributed equally to this research.

<sup>†</sup>Both authors contributed equally to this research.

---

Authors' Contact Information: Maria Teleki, Texas A&M University, College Station, U.S.A., [mariateleki@tamu.edu](mailto:mariateleki@tamu.edu); Anna Seo Gyeong Choi, Cornell University, Ithaca, U.S.A., [sc2359@cornell.edu](mailto:sc2359@cornell.edu); Anne Duray, AAAS Science and Technology Policy Fellow, Washington, D.C., U.S.A., [anne.m.duray@gmail.com](mailto:anne.m.duray@gmail.com); Haoran Liu, Texas A&M University, College Station, U.S.A., [liuhr99@tamu.edu](mailto:liuhr99@tamu.edu); Junyan Zhang, Independent, College Station, U.S.A., [zhangjunyan0225@gmail.com](mailto:zhangjunyan0225@gmail.com); Xiangjue Dong, Texas A&M University, College Station, U.S.A., [xj.dong@tamu.edu](mailto:xj.dong@tamu.edu); Dilma Da Silva, Texas A&M University, College Station, U.S.A., [dilma@tamu.edu](mailto:dilma@tamu.edu); Allison Koenecke, Cornell Tech, New York, U.S.A., [koenecke@cornell.edu](mailto:koenecke@cornell.edu); James Caverlee, Texas A&M University, College Station, U.S.A., [caverlee@tamu.edu](mailto:caverlee@tamu.edu).



This work is licensed under a Creative Commons Attribution 4.0 International License.

FAcCT '26, Montreal, QC, Canada

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2596-8/2026/06

<https://doi.org/10.1145/3805689.3812328>

**ACM Reference Format:**

Maria Teleki, Anna Seo Gyeong Choi, Anne Duray, Haoran Liu, Junyan Zhang, Xiangjue Dong, Dilma Da Silva, Allison Koenecke, and James Caverlee. 2026. The Due Process Deficit: Auditing AI Governance in U.S. Higher Education. In *The 2026 ACM Conference on Fairness, Accountability, and Transparency (FAccT '26)*, June 25–28, 2026, Montreal, QC, Canada. ACM, New York, NY, USA, 41 pages. <https://doi.org/10.1145/3805689.3812328>

**1 Introduction**

U.S. universities have met Artificial Intelligence (AI) with policy: revised honor codes, AI use guidelines, and detection tools to enforce them. What has not kept pace is the infrastructure needed to ensure that affected parties can understand, contest, and seek redress for decisions impacting them. This is a failure of *procedural justice* [39, 42]. A student flagged for academic misconduct by an AI detection tool [71] – with no published guidance on who can review the determination, what recourse exists, or whether any human is required to weigh in – has policy applied to them but may have no governance protecting them. When institutions do not provide this fundamental protection, they risk their own legitimacy [11, 39, 42].

Universities exercise significant power over students and faculty,<sup>1</sup> and AI-related governance decisions carry material consequences. A student accused of academic misconduct on the basis of an AI detection tool faces potential sanctions that can end a career [62]. Yet, as shown by our results, most institutions publish no guidance on how such determinations are made, whether human review is required, or what recourse is available. Faculty, meanwhile, are frequently expected to make consequential decisions about AI use – whether and how to permit it, how to detect misuse, how to adjudicate ambiguous cases – without adequate institutional support, resources, or training. And the policies that govern both students and faculty are rarely shaped by structured input from either group: mechanisms for participation, feedback, and revision are among the least developed components of institutional AI governance. Due process guarantees are legally protected in public universities under the Fourteenth Amendment, and both public and private institutions are bound by their own published procedures [62] – but published procedures must exist before they can be followed, enforced, or contested.

Public legibility, as we use the term, refers to the extent to which governance mechanisms – policies, procedures, feedback mechanisms, and recourse pathways – are findable, interpretable, and actionable by affected parties operating under realistic time and access constraints. Legibility is not sufficient for justice, but it is necessary [4, 12]. The following example illustrates why:

*Consider two students enrolled in the same course: Student A attends office hours regularly, has informal access to faculty, and learns through this interpersonal network that the instructor permits the use of generative AI for brainstorming. Student A therefore uses AI with confidence. Student B works full-time, attends class remotely, and relies primarily on the syllabus and university-level policies to determine acceptable practice. In the absence of clear, publicly articulated guidance, Student B either avoids AI use entirely or uses it cautiously while facing the perceived risk of academic integrity sanctions.*

Although both students are subject to the same formal rules, they encounter materially different guidance in practice. This failure is not neutral: it systematically advantages students and faculty with greater institutional access, while disadvantaging those who rely on official documentation and published policy. Informal governance is therefore a structural source of inequity, independent of institutional intent. The consequences are compounding: students report fear of being falsely accused of AI-assisted cheating [25, 30, 30, 33, 70], yet those who avoid AI entirely risk falling behind in the competencies their fields increasingly demand [13, 28]. Each of the failures Student B encounters – unclear policy, the assumption of faculty-specific rules, absent feedback mechanisms,

<sup>1</sup>We use *faculty* as an umbrella term for all instructional staff. We note that this label encompasses roles with differing levels of authority, job security, and participation in institutional governance, which vary across universities.

and ungoverned detection practices – represents a distinct dimension of governance capacity. Together, they illustrate core dimensions of procedural accountability.

Prior empirical audits of university AI governance have focused on policy presence among elite or research-intensive samples [47, 50, 76, 80] (Table 2, Appendix B), examining primarily the pedagogical dimension: how universities communicate expectations about AI use in coursework. We extend this work by auditing a structurally diverse sample of 79 U.S. universities spanning R1, R2, and teaching-focused institutions, and by shifting the analytic frame from pedagogy to *procedural accountability* [11, 12] – asking not whether universities have AI policies, but whether those policies constitute legible, contestable structures through which affected parties can understand the rules, participate in their revision, and seek recourse when things go wrong. A focus on procedural accountability also has more holistic implications, as pedagogical questions and expectations surrounding AI may have variance between disciplines.

We present **ACAI-US79**, an institutional audit dataset capturing publicly articulated AI governance practices across 79 diverse U.S. universities, and **ACAI**, the Academic AI Capacity Index—an interpretable measure of governance capacity grounded in the public legibility of formal policies, resources, feedback mechanisms, and procedural safeguards. The audit evaluates four governance domains: **A. POLICY CLARITY**, **B. FACULTY SUPPORT**, **C. FEEDBACK MECHANISMS**, and **D. DETECTION TOOLS**. These domains are organized around procedural logics: *Can affected parties identify the rules? Do faculty have the institutional support to apply them consistently? Are there structured mechanisms for participation and revision? And are the uses of AI detection technology – among the highest-stakes procedural interventions universities deploy – governed by safeguards against error and abuse?*

Our findings reveal that governance capacity is consistently concentrated in rule articulation – policy documents and faculty-facing guidance – while the mechanisms that enable participation, feedback, and procedural constraint are systematically underdeveloped. This pattern holds even among institutions with high aggregate governance capacity. The gap is not primarily explained by resource constraints or institutional type: AI research intensity is weakly correlated with AI governance capacity, indicating that technical leadership does not translate into procedural accountability (Appendix E.2). The structural problem is one of institutional design.

We make the following contributions, which together constitute the first large-scale, reproducible audit of procedural accountability in U.S. university AI governance:

- (1) A Structurally Diverse Institutional Audit:** Unlike prior audits restricted to elite or R1-only samples [47, 50, 76, 80] (Table 2, Appendix B), we audit 79 U.S. universities spanning R1, R2, and teaching-focused institutions, surfacing accountability gaps that top-N sampling systematically obscures. We release the dataset, annotation schema, audit toolkit, and public website to support transparency, replication, and institutional self-assessment (available at <https://github.com/mariateleki/ACAI> and <https://acai-us79.org/>).
- (2) A Legal and Institutional Analytic Frame:** We introduce ACAI, an interpretable audit index grounded in procedural accountability rather than pedagogical guidance – measuring whether authority, review, and recourse mechanisms are publicly legible to affected parties, and operationalizing due process as an empirically auditable property of institutional governance.
- (3) A Finding That Governance Capacity is Only Weakly Correlated with Research Activity:** We compare AI governance capacity (ACAI) against AI research output (CSRankings), finding that technical leadership does not translate into publicly articulated accountability – a result with direct implications for how the AI research community understands the relationship between AI investment and institutional responsibility.
- (4) A Critical Evaluation of Automated Governance Auditing:** We evaluate whether LLMs can approximate human governance judgments, finding that automated approaches produce unstable and misleading institutional rankings – showing the continued necessity of human judgment in high-stakes audit contexts.

**(5) Actionable Recommendations for Universities:** We translate audit findings into concrete recommendations for strengthening institutional AI governance, with particular attention to feedback mechanisms, human review requirements, and procedural safeguards around AI detection tools.

## 2 Related Work

### 2.1 AI in Higher Education

AI has achieved rapid, widespread adoption in higher education. Within two years of ChatGPT’s release, over 80% of students were using AI academically [21], and a 2025 survey found that 85% had used AI for coursework in the past year [28]. Adoption extends to instructors as well: 65% of STEM academic scientists report using generative AI in teaching or research activities [5]. Yet institutional preparedness has not kept pace. The American Association of Colleges and Universities notes that 68% percent of faculty survey respondents indicated their institutions had not prepared them for AI-related teaching challenges, while 78% reported increased cheating since generative AI tools became widely available [61]. Students themselves recognize the gap: more than half want education on ethical AI use, and more than half want clearer, standardized policies governing when and how AI tools can be used [28].

Calls for participatory governance in higher education have emerged from educators, professional organizations, and policy bodies alike. The EngageAI Task Force’s “AI Bill of Rights for Educators” [24] frames participatory governance as a matter of professional rights, calling for faculty agency, advocacy, and institutional support. The 2025 AAUP report [58] documents widespread faculty concern over exclusion from AI-related decision-making, and goes further by recommending the formation of standing or ad hoc committees composed of faculty, staff, and students—“elected by their constituencies and charged with monitoring, evaluating, and reviewing [AI] procurement processes and policy.” A recent EDUCAUSE working paper [32] similarly calls for increased stakeholder engagement, arguing that governance frameworks informed by diverse stakeholder voices are essential to realizing AI’s potential while protecting against harm. Several institutions have convened formal advisory structures — Cornell through a cross-unit Administrative AI Task Force [22], Stanford through a provost-commissioned advisory committee [1], and Michigan through a high-level AI Research Committee [73] — to evaluate risks and develop institutional recommendations. Yet student involvement is rarely built into these structures [36]. Further, systematic evidence of the effectiveness or durability of such participatory mechanisms remains limited, and we find that structured pathways for student and faculty input into AI governance are among the least developed components of institutional practice.

Despite widespread faculty concern about AI-assisted plagiarism [58], the detection tools deployed in response are neither accurate nor reliable – further discussed in Appendix C. False positive rates vary widely across tools, and detection systems disproportionately misclassify the work of non-native English speakers and neurodivergent writers [43, 77]. At institutional scale, even a 1% false positive rate (for example) translates into thousands of wrongful misconduct accusations [71], as shown in Figure 4. Yet institutional guidance on how detection outputs should be interpreted, weighted, or contested remains largely absent [47, 76].

### 2.2 Auditing Institutional AI Governance: Ethics and Accountability

Bommasani et al. [10] offers the closest precedent to our AI governance audit: the Foundation Model Transparency Index (FMTI) audits the publicly legible properties of AI systems themselves, producing a structured, comparative ranking across developers. Just as the FMTI asks whether AI developers make their systems legible enough to support external accountability, ACAI asks whether universities make their governance structures legible that affected parties can invoke due process protections.

Prior work has examined how universities articulate responses to generative AI, primarily through analyses of publicly available policies and guidelines. Previous studies analyze AI governance across only elite universities

(Table 2, Appendix B) – “Big 10” [80], “Top 100” [76], “Top 50” [50], or the 116 R1 universities in the U.S. [47] – focusing only on the pedagogical dimension: how universities communicate expectations about AI use in coursework. We extend this line of work in two directions: first, we audit a structurally diverse sample of 79 U.S. universities spanning R1, R2, and teaching/liberal arts institutions; and second, by shifting the analytic frame from pedagogy to *procedural accountability* [11, 12], asking whether authority, review, and recourse mechanisms are publicly legible to affected parties.

This framing is grounded in a sociotechnical understanding of accountability [38, 66, 79] (Appendix E.1): evaluating whether AI use in universities is accountable requires evaluating the institutional arrangements through which AI is authorized, constrained, and contested – not just the technical properties of the systems deployed. Algorithmic impact assessment scholarship reinforces this point. Metcalf et al. [49] and Selbst [65] document how opacity in algorithmic systems enables harm to go unaddressed. Ananny and Crawford [4] argue that “transparency alone cannot create accountable systems” – legibility is necessary but insufficient without mechanisms for participation and contestation.

These concerns are not merely theoretical: when governance structures lack public legibility, affected parties have no structured means of shaping or challenging the policies that govern AI use against them [68, 82]. This leads directly to ethical concerns (stemming from power imbalances), which are exacerbated when considering the compounding disproportionate harms to marginalized groups of students and faculty who hold less power within the university ecosystem. Our audit aims to ameliorate these potential harms through transparent, systemic measurements of university AI governance, which can in turn facilitate regulatory oversight or community-based advisory committees to mitigate the disparate impacts potentially arising from status quo, black-box AI governance.

### 2.3 Legal Grounding

ACAI is grounded in the legal tradition of algorithmic due process audits [41]. Universities are a setting where the *due process deficit* risk is acute: AI detection tools carry material consequences for students – including “financial loss and a stigma that may preclude admission to another university or access to a career” [62] – under governance frameworks that often lack the procedural specificity needed to support meaningful review. Public universities bear constitutional obligations under the Fourteenth Amendment; both public and private institutions are bound by their own published procedures and subject to judicial review [62]. Publicly legible governance is the precondition for these protections to function in practice. Kinchin [42] identifies the theoretical gap ACAI addresses: while algorithmic fairness has focused on distributive justice, a *procedural gap* persists in consent, transparency, and the right to be heard. Meaningful accountability requires information, explanation, and the possibility of consequences [11, 12]; contestation requires institutional infrastructure that affected parties can locate and invoke [39]. ACAI measures whether that infrastructure exists – a requirement for any meaningful right to contest.

## 3 An Audit of Institutional Capacity

We first clarify what we mean by institutional capacity and why public legibility provides an appropriate object of empirical audit. We conceptualize AI governance in higher education as a question of institutional capacity and accountability, not technical sophistication or ethical aspiration. We treat universities as powerful organizational actors that structure how AI systems are authorized, constrained, and contested through policies, procedures, and oversight mechanisms. Adopting an audit perspective, we examine the public legibility of these governance arrangements: what is made visible, enforceable, and contestable to students, faculty, and other affected parties. This framing motivates our audit design and ACAI as a tool for evaluating how governance capacity is institutionalized across universities.

You will review publicly available web pages for <UNIVERSITY> to determine whether specific AI-related policy statements are addressed by the institution.

Use ONLY the links provided below and any pages, sections, PDFs, or subpages that are directly reachable by clicking links on those pages (e.g., menus, internal links, or document links). Do NOT use external search engines or sources found from outside this list: <LINKS>.

Evaluate each statement independently. Spend no more than 5 minutes per statement.

**For each statement:**

**1. Select exactly one classification:**

- *Present/Yes* – A clear statement directly addressing the item is found on an institutional page within 5 minutes.
- *Partial/Implicit/Somewhat* – The item is mentioned or implied, but key details are missing.
- *Absent/No* – You reasonably searched the allowed sources and did not find relevant content.
- *Unclear or Took Longer Than 5 Minutes* – Navigation difficulty, vague language, or time limits prevented a confident decision.
- *Conflicting Information* – Different institutional sources provide contradictory guidance for the same item.

**2. Provide the most relevant URL(s)** from the allowed sources that support your selection. If you selected Absent or Unclear, provide the main page(s) you checked.

**Statements Organized by Governance Domain:**

**A. POLICY CLARITY** – *Policies defining institutional expectations, terminology, and academic integrity adaptations.*

A1. The university defines “AI use,” “AI assistance,” or “AI-generated content.”

A2. The university defines standards for citing AI-generated material.

**B. FACULTY SUPPORT** – *Resources that enable faculty to integrate, regulate, or teach with AI.*

B1. The university provides guidance, training, or resources for faculty on AI-related teaching practices.

B2. Official examples of appropriate and/or prohibited AI use are provided (e.g. example AI use cases, example prompts).

B3. A faculty committee or group focused on teaching and learning about AI exists.

B4. Faculty are offered syllabus language examples (e.g. use AI/don’t use AI/selectively use AI).

**C. FEEDBACK MECHANISMS** – *Mechanisms through which universities gather input, revise policies, and communicate decisions.*

C1. A faculty committee or advisory group focused on university AI policy or governance exists.

C2. A student committee or advisory group focused on university AI policy or governance exists.

C3. The university publishes AI policy update logs or explains revisions.

**D. DETECTION TOOLS** – *Institutional stance toward AI detection technologies.*

D1. The university restricts, discourages, or warns against the use of AI detection tools.

D2. Student misconduct determinations require human review and cannot be based solely on AI detection tools.

Fig. 1. **Annotation Instructions for ACAI Calculation, Organized By Governance Domain:** ACAI calculation details are provided in §3.2, governance domains are detailed in §3.2.1, and additional annotation details are in Appendix F.3.

### 3.1 ACAI-US79: A Benchmarking Dataset

Our annotation protocol – shown in Figure 1 – operationalizes institutional capacity as publicly legible, time-bounded, and normatively specific governance, rather than as the mere presence of AI-related content. Capacity is understood as what institutions make visible and actionable to affected stakeholders within reasonable time and effort constraints. The empirical focus on 79 U.S. universities is motivated by the country’s prominent position in global AI capacity indices [26, 48, 51, 55], detailed in Appendix E.2.

**University Selection.** The universities selected for the ACAI-US79 dataset and their corresponding attributes are shown in Table 1. The dataset was constructed using a purposive, diversity-oriented selection strategy designed to surface variation in how AI governance is institutionalized across U.S. higher education. The selection emphasizes institutional heterogeneity along dimensions that shape authority, accountability, and public legibility.

Within each U.S. Census region [74] – **South**, **West**, **Midwest**, and **Northeast** – we sought to include institutions spanning research intensity and institutional organization. Specifically, for each region we targeted approximately four institutions in each of the following categories: **Public Research** universities (**R1** or **R2**),

**Private Research** universities ( **R1** or **R2** ), and **Teaching/Liberal Arts** colleges, categorized according to the Carnegie Classifications [3]. Candidate institutions within each cell were identified using ChatGPT-assisted search (e.g., “R1 universities South region”) followed by manual verification against Carnegie Classification records. The choice of four institutions per category per region was a pragmatic design decision. Institution size (measured as number of students) was not used as an explicit stratification variable, thus the size distribution in ACAI-US79 emerges from the selection process and size is therefore treated analytically as a contextual attribute. We split size by tertiles into **Small**, **Medium**, and **Large** for comparison. Overall, this structure was intended to capture *differences in organizational incentives* that plausibly affect how AI governance is articulated.

**Link Retrieval.** For each university included in the audit, we systematically collected institutionally authoritative documents and official web links<sup>2</sup> across seven categories, with descriptions below and examples shown in Appendix F.2. These categories correspond to distinct organizational surfaces through which governance is commonly articulated [80], initialized by a manual author review of a small set of university policies, and then iterated on during our search process which is detailed in Appendix F.

The seven identified categories serve as search lenses: (1) University Policies or Guidelines: University-level policies addressing technology use, data governance, research ethics, or academic administration; (2) Center for Teaching & Learning: Centrally maintained guidance for faculty on pedagogical use of AI, syllabus adaptation, and instructional support; (3) AI Institute/Initiative/Center: Institutionally recognized units focused on AI, digital ethics, or related governance-relevant coordination; (4) Library Guide: Public-facing library resources addressing generative AI, citation practices, or responsible research use; (5) Academic Integrity: Policies or guidance governing the relationship between AI, authorship, plagiarism, and assessment; (6) AI Committee: Formal committees or task forces charged with evaluating or coordinating institutional responses to AI; and (7) Other Relevant Links: Additional institutionally authoritative materials relevant to AI governance.

The audit does not assume that governance capacity is limited to legally binding rules. Instead, we operationalize capacity as institutionalization: the presence of standing reference points – such as policies, offices, committees, or officially maintained resources – that persist over time and orient behavior by establishing expectations about authority, coordination, and acceptable practice. They represent *ongoing points of reference that a student, instructor, administrator, member of the press, or member of the public could reasonably consult* to understand how AI is governed at the institution. In contrast, transient communications such as news articles, announcements, or blog posts were excluded.<sup>3</sup> While such materials may signal institutional intent or activity, they do not establish durable roles, procedures, or accountability structures. Including them would collapse institutionalization into communication and systematically overstate governance capacity. Importantly, the links collected through this process (detailed in Appendix F) are treated as candidate surfaces of governance, not as governance determinations themselves. As shown in Figure 1, annotators were instructed to use these materials as evidence when evaluating whether specific governance statements were present, partial, or absent.

**Human Annotation via Prolific.** Annotation was conducted by paid human annotators recruited via the Prolific<sup>4</sup> platform, following the annotation instructions shown in Figure 1. Annotators labeled only publicly available institutional materials under standard compensated microtask conditions, and no PII or sensitive data were collected.<sup>5</sup> Tasks were estimated to take approximately 30 minutes and were compensated at \$12.00/hour, consistent with Prolific’s recommended compensation for U.S.-based participants. Each university was independently annotated by three annotators, administered via Google Forms. Participation was restricted to U.S.-based,

<sup>2</sup>Collection of links and annotation occurred from 12/8-12/30/25.

<sup>3</sup>Prior studies [47, 76, 80] can be referenced for details on AI governance via these communication types.

<sup>4</sup>Prolific is a crowdworker platform commonly used for academic research and data annotation tasks: <https://www.prolific.com/data-annotation>

<sup>5</sup>We provide the privacy policy in Figure 9.

English-fluent annotators using desktop or laptop devices, enforced through Prolific’s prescreening tools. To preserve the audit’s focus on public legibility, annotators were instructed not to use external search engines or prior knowledge, and each item was evaluated under a strict time limit. To support data quality, we enabled Prolific’s “reject exceptionally fast submissions” safeguard, which automatically flags submissions completed at implausibly short durations relative to the estimated task time, helping to filter inattentive or agent-driven responses. To assess alignment between expert and Prolific annotations, two computational PhD students independently annotated a subset of 10 universities using the same audit protocol. Agreement between the expert and aggregate Prolific annotations was moderate to strong ( $r = 0.55$ ,  $\rho = 0.53$ ,  $\tau = 0.48$ ), with Prolific scores exhibiting a small positive bias towards higher scores (as defined in §3.2, +10.45). Further, we compare our human annotated results to an LLM-driven audit, detailed in §3.4.

### 3.2 ACAI: The Academic AI Capacity Index

To enable systematic comparison of governance capacity across institutions, we construct **ACAI**, the **A**cademic **A**I Capacity **I**ndex. While the index necessarily collapses nuances present in more detailed datasets, ACAI serves as a first step toward understanding the AI governance landscape of U.S. universities. ACAI aggregates governance indicators across four domains (Figure 1) into a single interpretable score for university  $u$ :

$$\text{ACAI}_u = 100 \cdot \frac{\sum_{d \in \{A,B,C,D\}} w_d \left( \frac{1}{J_d} \sum_{j=1}^{J_d} \frac{1}{k} \sum_{i=1}^k I_{d,u,i,j} \right)}{\sum_{d \in \{A,B,C,D\}} w_d}$$

where  $I_{d,u,i,j} \in \{0, 0.5, 1\}$  is the score assigned by annotator  $i$  to indicator item  $j$  within governance domain  $d$  – where  $d \in \{A,B,C,D\}$ : **A. POLICY CLARITY**, **B. FACULTY SUPPORT**, **C. FEEDBACK MECHANISMS**, **D. DETECTION TOOLS**, and  $J_d$  is the number of indicators in governance domain  $d$ . Scores are coded as 1.0 (*Present/Yes*), 0.5 (*Partial/Implicit/Somewhat*), and 0.0 (*Absent/No, Unclear or Took Longer Than 5 Minutes, or Conflicting Information*). Domain scores are computed as the mean across indicators and annotators, and then aggregated via a weighted sum and scaled to  $[0, 100]$ . Fractional ranks (e.g., 2.5) arise from ties in the underlying ACAI scores, broken using average tie-ranking. Higher values indicate greater publicly articulated AI governance capacity. Weighting coefficients,  $w_d$ , encode normative priorities about which governance functions matter most; here we use the **indicator-weighted baseline** ( $w_A=2$ ,  $w_B=4$ ,  $w_C=3$ ,  $w_D=2$ ). To ensure findings are not artifacts of these choices, we assess robustness under alternative weighting schemes, varying  $1 \leq w_d \leq 4$  (see Finding 5 and Appendix H).

**A note on aggregation:** We intentionally aggregate indicator scores by averaging across annotators rather than resolving disagreement through expert adjudication or majority vote. This audit principle reflects the audit’s focus on public legibility rather than institutional intent, internal consistency, or expert interpretation. ACAI is explicitly not designed to capture how governance materials might be interpreted by legal counsel, administrators, or domain experts, but how they are encountered by external readers operating under realistic time and access constraints. When multiple annotators reviewing the same publicly available materials arrive at different judgments about whether a governance mechanism is present, partial, or absent, that variation is treated as an empirical signal of ambiguity in the underlying institutional artifacts. Hence, annotator disagreement is evidence about how clearly governance is articulated [2]. Averaging across multiple annotations aligns with the audit framing of ACAI as a measure of publicly articulated governance capacity: governance that requires expert interpretation or insider knowledge to interpret functions as weaker governance in practice, regardless of internal deliberation or intent. ACAI thus measures how governance is *encountered* by external readers under realistic constraints, not how it might be interpreted by insiders. Because the index relies on publicly available materials, it should be interpreted as a lower bound on institutional capacity.

**3.2.1 Governance Domains.** We now turn to the four governance domains covered in our study; explicit annotation directions are provided in Table 1, and specific examples are provided in Table 6 (Appendix G):

**A. POLICY CLARITY.** This domain captures the extent to which universities publicly articulate clear, institution-level expectations regarding AI use in academic contexts, following a long legal tradition establishing the importance of clear definitions to support stable legal interpretation [31, 44]. Policy clarity includes the definition of key terms (e.g., “AI use,” “AI assistance,” or “AI-generated content”), guidance on attribution or citation of AI-generated material, and the adaptation of existing academic integrity frameworks to account for AI-mediated authorship. Prior work has shown heterogeneity in how universities define and communicate AI-related expectations, with many institutions relying on vague or decentralized guidance [47, 76]. In the absence of clear, publicly legible policy language, responsibility for interpreting acceptable AI use is often shifted to individual faculty or students, increasing the risk of inconsistent enforcement and inequitable outcomes [58]. Policy clarity therefore functions as a foundational component of institutional governance capacity, establishing shared reference points for authority, compliance, and contestation.

**B. FACULTY SUPPORT.** This domain reflects the extent to which universities provide institutionally maintained resources that enable faculty to engage with AI in teaching and assessment in informed and supported ways – ultimately to support student learning [46]. Faculty support includes guidance on pedagogical uses of AI, training opportunities, example use cases or prohibitions, model syllabus language, AI professional development opportunities, and the presence of faculty-focused committees or working groups concerned with AI and teaching [47, 58, 76]. Faculty are frequently expected to make consequential decisions about AI use – such as whether and how to permit AI in coursework – without adequate institutional support or coordination [47, 58]. The presence of formal faculty support mechanisms signals that responsibility for AI governance is not delegated entirely to individual faculty, but is instead recognized also as an institutional obligation requiring shared infrastructure and expertise. Importantly, faculty support directly impacts students: clear guidance helps faculty communicate consistent expectations to students, while faculty committees and resources provide the foundation for student-facing policies articulated in the other domains.

**C. FEEDBACK MECHANISMS.** This domain reflects concerns articulated in prior survey research and aligns with recommendations from the July 2025 AAUP report *Artificial Intelligence and Academic Professions* [58]. The report draws on approximately 500 survey responses from AAUP members regarding their experiences with AI and other educational technologies. A central finding was a widespread concern (reported by 71% of respondents) about the disconnect between administrative decision-making on AI policy and meaningful faculty and student input. In response, the report recommends the adoption of “meaningful shared governance policies and practices,” including committees composed of faculty, staff, and/or students, as well as increased transparency around AI-related decisions and policy changes. More broadly, scholarship on algorithmic governance emphasizes that accountability requires institutionalized mechanisms for participation, feedback, and revision over time. Absent such mechanisms, governance frameworks risk functioning as static or symbolic commitments rather than durable, contestable structures of authority [49, 65]. In Table 6, we show an example of such a committee. For an extended discussion on this domain, see Appendix E.1.

**D. DETECTION TOOLS.** This domain captures institutional stances toward AI detection technologies used in academic integrity enforcement [71, 81]. Indicators in this domain assess whether universities restrict, discourage, or explicitly govern the use of AI detection tools, and whether procedural safeguards – such as requirements for human review – are articulated [42]. Prior research – shown in Table 4 – has documented significant technical limitations and bias in AI detection systems, as well as their potential to produce false positives with serious consequences for students [43, 76]. Despite these risks, institutional guidance on detection tools is often limited, ambiguous, or silent on procedural constraints [47], potentially algorithmically shortcutting due process<sup>6</sup>

<sup>6</sup>Khattak [41] on the U.S. context: “As artificial intelligence systems increasingly assist decisionmaking in judicial and administrative processes, courts and administrative agencies face mounting pressure to merge innovation with legal tradition. These technologies are often praised for their efficiency. However, when the mechanisms by which they operate are impenetrable, they threaten to infringe upon core due process protections. The

| ACAI Rank | ACAI Score | CSRankings <sub>AI</sub> Rank | Institution                                 | Type             | Research Activity | Region    | Size   | LLM <sub><math>\tau=1</math></sub> Rank | ACAI – LLM <sub><math>\tau=1</math></sub> |
|-----------|------------|-------------------------------|---------------------------------------------|------------------|-------------------|-----------|--------|-----------------------------------------|-------------------------------------------|
| 1         | 81.82      | 105                           | University of New Hampshire                 | Public Research  | R1                | Northeast | Medium | 11                                      | –10                                       |
| 2.5       | 80.30      | 96                            | Portland State University                   | Public Research  | R2                | West      | Large  | 24                                      | –21.5                                     |
| 2.5       | 80.30      | 8                             | Stanford University                         | Private Research | R1                | West      | Medium | 2.5                                     | 0                                         |
| 4         | 77.27      | 14                            | University of Texas at Austin               | Public Research  | R1                | South     | Large  | 24                                      | –20                                       |
| 5         | 75.76      | 42                            | University of Notre Dame                    | Private Research | R1                | Midwest   | Medium | 77.5                                    | –72.5                                     |
| 6.5       | 74.24      | 137                           | Baylor University                           | Private Research | R1                | South     | Large  | 39.5                                    | –33                                       |
| 6.5       | 74.24      | 44                            | University at Buffalo                       | Public Research  | R1                | Northeast | Large  | 33.5                                    | –27                                       |
| 9.5       | 71.21      | un.                           | Rowan University                            | Public Research  | R2                | Northeast | Large  | 19                                      | –9.5                                      |
| 9.5       | 71.21      | 29                            | Stony Brook University                      | Public Research  | R1                | Northeast | Large  | 21.5                                    | –12                                       |
| 9.5       | 71.21      | 79                            | University of Florida                       | Public Research  | R1                | South     | Large  | 33.5                                    | –24                                       |
| 9.5       | 71.21      | 9                             | University of Michigan at Ann Arbor         | Public Research  | R1                | Midwest   | Large  | 11                                      | –1.5                                      |
| 14        | 69.70      | 91                            | Case Western Reserve University             | Private Research | R1                | Midwest   | Medium | 6.5                                     | +7.5                                      |
| 14        | 69.70      | un.                           | Lafayette College                           | Teaching/LA      | —                 | Northeast | Small  | 33.5                                    | –19.5                                     |
| 14        | 69.70      | un.                           | Lewis & Clark College                       | Teaching/LA      | —                 | West      | Small  | 50.5                                    | –36.5                                     |
| 14        | 69.70      | 27                            | Texas A&M University                        | Public Research  | R1                | South     | Large  | 54.5                                    | –40.5                                     |
| 14        | 69.70      | 7                             | University of California, Berkeley          | Public Research  | R1                | West      | Large  | 16.5                                    | –2.5                                      |
| 20        | 68.18      | 117                           | Brandeis University                         | Private Research | R1                | Northeast | Small  | 27.5                                    | –7.5                                      |
| 20        | 68.18      | un.                           | California State University, Long Beach     | Public Research  | R2                | West      | Large  | 14.5                                    | +5.5                                      |
| 20        | 68.18      | un.                           | Chapman University                          | Private Research | R2                | West      | Medium | 27.5                                    | –7.5                                      |
| 20        | 68.18      | 6                             | Cornell University                          | Private Research | R1                | Northeast | Medium | 8                                       | +12                                       |
| 20        | 68.18      | un.                           | Howard University                           | Private Research | R1                | South     | Medium | 39.5                                    | –19.5                                     |
| 20        | 68.18      | 169                           | Southern Methodist University               | Private Research | R1                | South     | Medium | 52.5                                    | –32.5                                     |
| 20        | 68.18      | 25                            | University of North Carolina at Chapel Hill | Public Research  | R1                | South     | Large  | 11                                      | +9                                        |
| 24.5      | 66.67      | 100                           | Syracuse University                         | Private Research | R1                | Northeast | Large  | 27.5                                    | –3                                        |
| 24.5      | 66.67      | 81                            | University of South Florida                 | Public Research  | R1                | South     | Large  | 19                                      | +5.5                                      |
| 27        | 65.15      | 37                            | The Ohio State University                   | Public Research  | R1                | Midwest   | Large  | 21.5                                    | +5.5                                      |
| 27        | 65.15      | 15                            | University of Southern California           | Private Research | R1                | West      | Large  | 38                                      | –11                                       |
| 27        | 65.15      | un.                           | University of Wyoming                       | Public Research  | R1                | West      | Medium | 50.5                                    | –23.5                                     |
| 32        | 63.64      | 29                            | Arizona State University                    | Public Research  | R1                | West      | Large  | 57.5                                    | –25.5                                     |
| 32        | 63.64      | 142                           | DePaul University                           | Private Research | R2                | Midwest   | Large  | 11                                      | +21                                       |
| 32        | 63.64      | 169                           | Florida Institute of Technology             | Private Research | R2                | South     | Small  | 33.5                                    | –1.5                                      |
| 32        | 63.64      | un.                           | Fordham University                          | Private Research | R2                | Northeast | Medium | 4.5                                     | +27.5                                     |
| 32        | 63.64      | un.                           | Mercer University                           | Private Research | R2                | South     | Medium | 11                                      | +21                                       |
| 32        | 63.64      | un.                           | Northern Illinois University                | Public Research  | R2                | Midwest   | Medium | 4.5                                     | +27.5                                     |
| 32        | 63.64      | un.                           | University of South Alabama                 | Public Research  | R2                | South     | Medium | 64                                      | –32                                       |
| 36        | 62.12      | un.                           | Illinois State University                   | Public Research  | R2                | Midwest   | Large  | 1                                       | +35                                       |
| 39        | 60.61      | 100                           | Binghamton University                       | Public Research  | R1                | Northeast | Medium | 33.5                                    | +5.5                                      |
| 39        | 60.61      | un.                           | Lake Forest College                         | Teaching/LA      | —                 | Midwest   | Small  | 66                                      | –27                                       |
| 39        | 60.61      | un.                           | Montclair State University                  | Public Research  | R2                | Northeast | Medium | 2.5                                     | +36.5                                     |
| 39        | 60.61      | un.                           | Pepperdine University                       | Private Research | R2                | West      | Medium | 33.5                                    | +5.5                                      |
| 39        | 60.61      | 46                            | University of Chicago                       | Private Research | R1                | Midwest   | Large  | 24                                      | +15                                       |
| 43.5      | 57.58      | un.                           | Carleton College                            | Teaching/LA      | —                 | Midwest   | Small  | 6.5                                     | +37                                       |
| 43.5      | 57.58      | 79                            | Iowa State University                       | Public Research  | R1                | Midwest   | Large  | 19                                      | +24.5                                     |
| 43.5      | 57.58      | un.                           | San José State University                   | Public Research  | R2                | West      | Large  | 33.5                                    | +10                                       |
| 43.5      | 57.58      | 16                            | University of Washington-Seattle            | Public Research  | R1                | West      | Large  | 48.5                                    | –5                                        |
| 48        | 56.06      | un.                           | Colby College                               | Teaching/LA      | —                 | Northeast | Small  | 60                                      | –12                                       |
| 48        | 56.06      | un.                           | Reed College                                | Teaching/LA      | —                 | West      | Small  | 41                                      | +7                                        |
| 48        | 56.06      | un.                           | Saint Louis University                      | Private Research | R1                | Midwest   | Medium | 27.5                                    | +20.5                                     |
| 48        | 56.06      | un.                           | Skidmore College                            | Teaching/LA      | —                 | Northeast | Small  | 43                                      | +5                                        |
| 48        | 56.06      | un.                           | Southern University and A & M College       | Public Research  | R2                | South     | Small  | 77.5                                    | –29.5                                     |
| 52.5      | 54.55      | un.                           | Davidson College                            | Teaching/LA      | —                 | South     | Small  | 63                                      | –10.5                                     |
| 52.5      | 54.55      | un.                           | Southern Wesleyan University                | Teaching/LA      | —                 | South     | Small  | 70                                      | –17.5                                     |
| 52.5      | 54.55      | un.                           | University of Colorado Colorado Springs     | Public Research  | R2                | West      | Medium | 60                                      | –7.5                                      |
| 52.5      | 54.55      | 2                             | University of Illinois Urbana-Champaign     | Public Research  | R1                | Midwest   | Large  | 14.5                                    | +38                                       |
| 55        | 53.03      | 142                           | Wichita State University                    | Public Research  | R2                | Midwest   | Medium | 43                                      | +12                                       |
| 57        | 51.52      | un.                           | Georgia Southern University                 | Public Research  | R2                | South     | Large  | 68.5                                    | –11.5                                     |
| 57        | 51.52      | un.                           | University of Denver                        | Private Research | R1                | West      | Medium | 46                                      | +11                                       |
| 57        | 51.52      | un.                           | Wofford College                             | Teaching/LA      | —                 | South     | Small  | 43                                      | +14                                       |
| 59        | 50.00      | un.                           | Westminster University                      | Teaching/LA      | —                 | West      | Small  | 67                                      | –8                                        |
| 62        | 48.48      | un.                           | Abilene Christian University                | Private Research | R2                | South     | Small  | 77.5                                    | –15.5                                     |
| 62        | 48.48      | un.                           | Ball State University                       | Public Research  | R2                | Midwest   | Large  | 52.5                                    | +9.5                                      |
| 62        | 48.48      | 49                            | Brown University                            | Private Research | R1                | Northeast | Medium | 56                                      | +6                                        |
| 62        | 48.48      | un.                           | Clark University                            | Private Research | R2                | Northeast | Small  | 73                                      | –11                                       |
| 62        | 48.48      | un.                           | Rhodes College                              | Teaching/LA      | —                 | South     | Small  | 46                                      | +16                                       |
| 66        | 46.97      | 100                           | Illinois Institute of Technology            | Private Research | R2                | Midwest   | Small  | 48.5                                    | +17.5                                     |
| 66        | 46.97      | 142                           | Nova Southeastern University                | Private Research | R1                | South     | Large  | 62                                      | +4                                        |
| 66        | 46.97      | un.                           | Wesleyan University                         | Teaching/LA      | —                 | Northeast | Small  | 54.5                                    | +11.5                                     |
| 68        | 45.45      | un.                           | Kean University                             | Public Research  | R2                | Northeast | Medium | 65                                      | +3                                        |
| 70        | 43.94      | 71                            | California Institute of Technology          | Private Research | R1                | West      | Small  | 46                                      | +24                                       |
| 70        | 43.94      | un.                           | Occidental College                          | Teaching/LA      | —                 | West      | Small  | 75                                      | –5                                        |
| 70        | 43.94      | 61                            | Stevens Institute of Technology             | Private Research | R2                | Northeast | Small  | 60                                      | +10                                       |
| 72        | 42.42      | un.                           | Long Island University                      | Private Research | R2                | Northeast | Medium | 71.5                                    | +0.5                                      |
| 73        | 40.91      | un.                           | Marquette University                        | Private Research | R2                | Midwest   | Medium | 16.5                                    | +56.5                                     |
| 74        | 39.39      | un.                           | Beloit College                              | Teaching/LA      | —                 | Midwest   | Small  | 77.5                                    | –3.5                                      |
| 75        | 34.85      | un.                           | Grinnell College                            | Teaching/LA      | —                 | Midwest   | Small  | 33.5                                    | +41.5                                     |
| 76        | 31.82      | un.                           | Jackson State University                    | Public Research  | R2                | South     | Small  | 57.5                                    | +18.5                                     |
| 77        | 28.79      | un.                           | Clark Atlanta University                    | Private Research | R2                | South     | Small  | 71.5                                    | +5.5                                      |
| 78.5      | 27.27      | un.                           | Creighton University                        | Private Research | R2                | Midwest   | Medium | 68.5                                    | +10                                       |
| 78.5      | 27.27      | un.                           | University of Massachusetts at Dartmouth    | Public Research  | R2                | Northeast | Small  | 74                                      | +4.5                                      |

Table 1. **Institutional Rankings for ACAI-US79 with ACAI and CSRankings<sub>AI</sub> indicator-weighted variant, detailed in Finding F5 and Appendix H):** Details for ACAI calculations in §3; *Region* is classified based on the U.S. Census [74]; *Type* and *Research Activity* are classified based on the Carnegie Classifications [3]; *Size* (# students) is split based on tertile buckets.  $LLM_{\tau=1}$  is the LLM-generated ACAI rank at temperature  $\tau = 1.0$ ;  $ACAI - LLM_{\tau=1}$  reports the directional rank gap (green = LLM overestimated, red = LLM underestimated, grey = exact agreement). Within tied ACAI scores, institutions are listed alphabetically by name.




Fig. 2. **AI governance capacity scales with institutional type, research activity, and size, but its composition is skewed at every level: Policy Clarity (A) and Faculty Support (B) outscore Feedback Mechanisms (C) and Detection Tools (D), leaving participation and due process systematically undersupported. (>F1, F2). 95% CIs shown.**

protections [41, 62], and with little-to-no contestability [45]. Publicly articulated governance in this domain is therefore critical for clarifying the role of detection tools in decision-making, delineating authority between automated systems and human judgment, and protecting due process in academic misconduct determinations. For an extended discussion on this subdomain, see Appendix C.

### 3.3 Approximating AI Research Activity with CSRankings<sub>AI</sub>

To contextualize institutional AI governance capacity relative to AI research activity, we used the rankings provided by CSRankings, a widely used, publicly available ranking of computer science research output shown in Figure 11. CSRankings aggregates publication counts across DBLP computer science venues and allows filtering by research area and time period; for details on the specific configuration used in this study, see Appendix I. As shown in Table 1, we report the CSRankings<sub>AI</sub> rank of each ACAI-US79 university to facilitate our comparison of AI governance capacity and AI research activity, and examine correlations between these two rankings as shown in Figure 10. Notably, CSRankings only ranks research-active institutions by publication output; Teaching/Liberal Arts colleges and institutions without substantial computer science research activity do not appear in CSRankings and are therefore excluded from the correlation analysis in Figure 10.

### 3.4 LLM Study

Because universities increasingly rely on AI systems to assess, classify, and enforce academic norms [16, 27, 37, 63], we conducted an LLM study with gpt-5.2-2025-12-11 to examine whether similar systems can meaningfully evaluate institutional AI governance: we conducted an LLM study aligned with the ACAI-US79 audit framework. The prompt used is shown in Figure 12; we perform three independent runs of for each sampling temperature

*Constitution guarantees that individuals be informed of decisions affecting their rights, and to have a fair opportunity to contest those decisions in a meaningful way. When unclear algorithms replace human judgment, those guarantees are at risk of becoming procedural only by name.”*

$\tau \in \{0.5, 1.0, 1.5\}$ . For each university-statement pair, the model was required to return a single categorical score and supporting URLs in a strictly validated JSON schema.

## 4 Findings

We use exploratory comparisons at both the aggregate and subgroup levels<sup>7</sup> to examine how AI governance capacity is differentially institutionalized across higher-education contexts. Contrasts in ACAI scores across institutional type, research activity, region, and size are used as a descriptive lens on organizational variation.

▷ **F1: AI governance capacity varies with organizational structure.** As shown in Figure 2, **Public Research** universities tend to occupy higher positions in the ACAI distribution than **Private Research** or **Teaching/Liberal Arts** universities; larger institutions (**Large** and **Medium**) tend to exhibit higher ACAI scores than **Small** ones; **R1** schools tend to show higher ACAI scores than **R2** and unclassified (–) schools; and schools in the **Western** region tend to show slightly higher ACAI scores. These differences align with known variation in institutional oversight arrangements and coordination demands: public universities commonly operate under statutory or regulatory accountability frameworks and maintain centralized administrative infrastructures, which are visible in the form of institution-level policies, guidance pages, and standing committees [8, 62, 78]. In contrast, institutions characterized by more decentralized organizational structures often rely on localized practices, which may be less consistently reflected in publicly accessible governance artifacts. Additionally, large institutions typically coordinate governance across a greater number of academic units, faculty, and students, which is reflected in the presence of centrally maintained and publicly legible reference points. Overall, even among the highest-capacity institutions, ACAI scores reveal substantial room for improvement. For example, UC Berkeley – the highest-scoring institution among Large, Public Research, R1 universities in the West – achieves an ACAI score of 69.70.

Importantly, ACAI does not assess research quality, ethical commitments, or internal decision-making processes. It captures whether AI-related governance mechanisms are publicly articulated and institutionally maintained. From this perspective, lower ACAI scores should not be interpreted as evidence of weaker concern or expertise, but as indicative of different approaches to organizing and communicating governance. These findings highlight a structural tension: institutional arrangements that emphasize decentralization or flexibility may be less visible in public-facing governance artifacts, even when substantive internal practices are present. However, this lack of public legibility introduces a specific accountability risk: reliance on informal guidance or word-of-mouth governance differentially advantages actors with greater institutional access, while disadvantaging students and faculty who must rely on publicly accessible rules and procedures. Publicly articulated governance capacity thus functions as a necessary condition for procedural accountability, establishing a minimally equitable baseline that does not depend on insider knowledge.

▷ **F2: AI governance capacity is concentrated in policy articulation rather than participation or process.** Figure 2 reveals that across institutional types, sizes, research intensities, and regions, scores in **A. POLICY CLARITY** and **B. FACULTY SUPPORT** are systematically higher than those in **C. FEEDBACK MECHANISMS** and **D. DETECTION TOOLS**. This consistent structural pattern indicates that publicly articulated AI governance capacity is concentrated in domains oriented toward rule articulation and instructional guidance, rather than in mechanisms that enable participation, feedback, or procedural constraint.

Notably, this gap persists even among institutions with otherwise high aggregate ACAI scores. **Large**, **Public Research**, and **R1** universities – while exhibiting higher overall governance capacity – still show pronounced deficits in feedback and detection tool governance relative to policy articulation. This pattern

<sup>7</sup>These subgroup comparisons are not treated as confirmatory evidence: the audit sample is purposive rather than probabilistic, subgroup categories are administratively defined and analytically coarse, and multiple overlapping contrasts are examined without correction.




Fig. 3. **Directional rank gap,  $\text{ACAI} - \text{LLM}_{\tau=1}$ , sorted descending ( $>F4$ ).** As shown in Table 1, positive values indicate that the LLM overestimated governance capacity (e.g., Marquette:  $73 - 16.5 = +56.5$ ); negative values indicate underestimation (e.g., Notre Dame:  $5 - 77.5 = -72.5$ ). Lollipop color encodes the quartile of the absolute misranking  $|\Delta|$ : purple marks the top 25% (the most extreme misrankings,  $|\Delta| \geq 24$ ), blue the middle 50% ( $6.5 < |\Delta| < 24$ ), and yellow the bottom 25% (the smallest misrankings,  $|\Delta| \leq 6.5$ ). Correlations across temperatures are reported in Table 11. We notice multiple failure modes behind these misrankings, the most notable being that the LLM marked **all** retrieved content as *Unclear* (which the scoring pipeline maps to 0) (e.g., Notre Dame). We leave a more in-depth investigation to future work.

suggests that differences are primarily quantitative rather than qualitative: AI governance capacity scales with organizational resources, but its internal composition remains skewed toward static guidance rather than durable procedural safeguards. In this sense, higher-capacity institutions often extend the same governance model rather than adopting qualitatively different forms of participatory or process-oriented governance.

As shown in Figure 2, although AI detection tools are frequently referenced in academic integrity materials, explicit procedural guidance governing their use is rare. This aligns with the findings of Wang et al. [76], who found that while 57% of universities in their dataset mentioned common tools, none explicitly recommended their use. Institutions often fail to specify whether detection tools are advisory or determinative, how results should be interpreted, what safeguards exist against error, or what recourse is available to affected students. This gap creates ambiguity around authority and enforcement, and risks inconsistent or discretionary application in practice, as expanded upon in Appendix C, directly harming students, and creating the potential for an adversarial relationship between students and faculty which is counterproductive to a healthy learning environment.

▷ **F3: AI governance capacity is weakly correlated with AI research output.** AI governance capacity (ACAI) is weakly correlated with AI-specific research output ( $\text{CSRanks}_{AI}$ ): across the  $N = 35$  institutions that both ACAI and  $\text{CSRanks}_{AI}$  rank, the rank correlation is weak: Pearson  $r = 0.17$ , Spearman  $\rho = 0.19$  (95% bootstrap CI  $[-0.15, 0.47]$ ), Kendall  $\tau = 0.12$ . Further, as shown in Figure 10, the subgroup Spearman 95% CIs consistently overlap with 0; this dispersion indicates that AI technical leadership alone does not reliably translate into strong, visible governance practices. This is further demonstrated in Table 1: The highest ACAI score is achieved by the University of New Hampshire, followed by Portland State University and Stanford University. Two of the top three institutions are **Public Research** universities, including one classified as **R2**, and three within the top 20 are **Small** institutions, two of which are solely **Teaching/Liberal Arts** colleges, illustrating that high levels of publicly articulated governance capacity are observed across a range of institutional types and research classifications. Across the full sample, several **Public Research** / **R2** universities – such

as Rowan University and California State University, Long Beach – appear in higher positions in the ACAI distribution than many **Private Research / R1** universities with substantial AI research activity. Conversely, multiple **Private Research / R1** universities – such as Brown University and the California Institute of Technology – appear in the lower half of the ACAI distribution despite significant contributions to AI scholarship. These contrasts indicate that AI research intensity and AI governance capacity are orthogonal dimensions of institutional capability.

▷ **F4: LLMs only partially reproduce human governance judgments.** Table 11 (Appendix J) reports full correlation statistics between LLM-generated rankings and human rankings across temperatures; Figure 3 visualizes the per-institution rank gap at  $\tau = 1.0$ . Across all temperatures, LLM-generated rankings exhibit moderate to strong ordinal agreement with human judgments ( $\rho = 0.54 - 0.62$ ), but individual institutions are frequently and severely misranked, with mean absolute rank errors of 15–17 positions out of 79. The errors are not symmetric: institutions in the top quartile of the rank-gap distribution experience misrankings of up to 64–72 positions – nearly inverting their relative ordering – while the bottom quartile shows nontrivial but modest discrepancies of 0–8 positions. Taken together, these results suggest that while LLMs may approximate coarse aggregate patterns, they fail to reliably reproduce the fine-grained, interpretive distinctions required for institutional governance audits. Governance evaluation depends on contextual reading, procedural inference, and judgment under ambiguity – capacities that are not robustly captured by current LLM-based approaches. This shows the continued necessity of human judgment in governance evaluation.

▷ **F5: ACAI rankings are robust to weighting choices and individual annotators.** To assess whether ACAI rankings are artifacts of either normative aggregation choices or variation in annotator judgments, we conducted a series of robustness analyses varying both domain weights and annotator inclusion. We first evaluated four weighting schemes: an indicator-weighted baseline ( $w_A=2, w_B=4, w_C=3, w_D=2$ ), equal ( $w_A=1, w_B=1, w_C=1, w_D=1$ ), policy-heavy ( $w_A=1, w_B=1, w_C=2, w_D=2$ ), and teaching-heavy ( $w_A=1, w_B=2, w_C=1, w_D=1$ ), using percentile ranks. Rankings were highly stable across weighting schemes (Spearman  $\rho = 0.93-0.99$ , Pearson  $r = 0.93-0.99$ ). Mean maximum rank shifts were a modest 16.32 ranks, with an interquartile range of 15.00–22.5 ranks and a maximum of 46 ranks; details in Appendix H. We compute inter-annotator agreement: at the level of individual items, agreement is modest and variable (Krippendorff's  $\alpha = 0.24$ ), reflecting the interpretive and normative nature of governance assessment and the heterogeneity of institutional documentation. Average pairwise agreement shows a similar pattern (mean = 0.49), indicating systematic but incomplete convergence among annotators. When indicators are aggregated at the governance domain-level, agreement improves ( $\alpha = 0.30$ ), suggesting that higher-level governance constructs are more consistently interpretable than individual policy statements. We also evaluated robustness to individual annotators using leave-one-annotator-out recomputation of ACAI scores under the indicator-weighted scheme. Across all three exclusions, recomputed scores remained strongly correlated with the full-annotator index ( $\rho = 0.77-0.87, r = 0.83-0.87$ ), indicating that no single annotator systematically altered the relative ordering of high- or low-capacity institutions.

#### 4.1 University Policy Recommendations

Based on the audit findings, we outline a set of policy recommendations aimed at strengthening institutional AI governance capacity in higher education.

▷ **R1: Establish a centralized institution-level AI governance reference point (following ▷F1, F3).** Universities should maintain clearly identifiable, centrally managed points of reference – such as policies, standing guidance pages, or designated offices – that articulate how AI use is governed. Reliance on informal norms or dispersed documentation makes governance difficult to locate and unevenly accessible to students and faculty.

▷ **R2: Provide procedural clarity around academic integrity and AI detection tools (following ▷F1, F2).** Where AI detection tools or integrity enforcement mechanisms are referenced, institutions should articulate clear procedures governing their use, limits, appeal processes, and responsible parties. Absent such guidance, detection practices risk being experienced as opaque, discretionary, or punitive.

▷ **R3: Formalize feedback and revision mechanisms (following ▷F2).** Governance capacity is strengthened when institutions specify how AI-related policies are reviewed, updated, and contested over time. Standing committees, task forces, or revision timelines signal that governance is ongoing rather than static or symbolic.

## 5 Discussion

**Conclusion.** We present a large-scale institutional audit of publicly articulated AI governance in U.S. higher education, shifting attention from technical systems to the organizational infrastructures through which AI-related authority is exercised. Our comparative ranking, ACAI, reveals substantial unevenness: governance capacity concentrates in rule articulation over mechanisms for participation, feedback, or procedural safeguards. Higher ACAI rankings do not consistently align with AI research intensity, indicating that institutional accountability is shaped more by organizational design and incentives than by technical leadership. We publicly release ACAI-US79, the audit instrument, and an accompanying website at <http://acai-us79.org/> and code at <https://github.com/mariateleki/ACAI> that makes governance artifacts directly inspectable and enables future re-audits. By rendering AI governance structures visible and comparable, this work provides an empirical foundation for advancing accountable AI governance in higher education.

**Scope & Limitations.** First, our sample is limited to 79 universities, which, although diverse in size and mission, cannot capture the full range of institutional practices globally. Second, the analysis is USA-centric, reflecting the regulatory, cultural, and policy context of U.S. higher education, but may not generalize to other geographies. Third, our reliance on publicly available institutional data may omit informal practices or internal decision-making processes that shape outcomes but are not externally visible. Fourth, our audit design treats governance as either publicly legible or informal, but does not capture a middle ground: formal, institutionally maintained materials accessible to enrolled students and staff — such as intranet pages, learning management systems, or password-protected policy portals — that are not publicly available to external observers. Such materials may constitute genuine governance infrastructure, and their exclusion means ACAI should be interpreted as a lower bound on governance capacity.

**Future Directions.** Future work should audit how emerging forms of AI-mediated surveillance and labor automation jointly erode institutional governance capacity in higher education. Universities are rapidly deploying AI-based monitoring systems under the language of integrity and efficiency — from AI detection tools (see §3.2.1, Appendix C) to fully automated proctoring based on behavioral analytics like eye- and click-tracking [63]. Yet these systems operate within asymmetrical power relations that render consent effectively coercive [17, 20, 64]: students cannot meaningfully opt out without material penalty, and governance mechanisms lag behind technological adoption. Compounding this, universities are automating AI-mediated instruction and auditing while reducing human interpretive labor [18, 60], displacing the very actors — TAs, faculty, and staff — who translate policy into practice and provide critical feedback on institutional decisions. Fully automated verbal exams using voice AI [37] and LLM-as-a-judge grading [16, 27] extend these dynamics into new domains, raising further questions for the governance frameworks we have outlined.

### Generative AI Usage Statement

We responsibly used AI technologies (ChatGPT-5.0) in this paper to assist with search, and the styling and language of the writing, as well as code assistance.

## References

- [1] AI at Stanford Advisory Committee. 2024. *AI at Stanford: Report of the Advisory Committee*. Technical Report. Stanford University, Office of the Provost. <https://provost.stanford.edu/sites/g/files/sbiybj29391/files/media/file/ai-at-stanford-report.pdf>
- [2] Dalia Ali, Dora Zhao, Allison Koenecke, and Orestis Papakyriakopoulos. 2025. Operationalizing Pluralistic Values in Large Language Model Alignment Reveals Trade-offs in Safety, Inclusivity, and Model Behavior. *arXiv preprint arXiv:2511.14476* (2025).
- [3] American Council on Education. 2025. Carnegie Classification of Institutions of Higher Education®. <https://carnegieclassifications.acenet.edu/>.
- [4] Mike Ananny and Kate Crawford. 2018. Seeing without knowing: Limitations of the transparency ideal and its application to algorithmic accountability. *new media & society* 20, 3 (2018), 973–989.
- [5] Wenceslao Arroyo-Machado, Jinghuan Ma, Tipeng Chen, Timothy P Johnson, Shaika Islam, Lesley Michalegko, and Eric Welch. 2025. Generative AI and academic scientists in US universities: Perception, experience, and adoption intentions. *PloS one* 20, 8 (2025), e0330416.
- [6] Clare Baek, Tamara Tate, and Mark Warschauer. 2024. “ChatGPT seems too good to be true”: College students’ use and perceptions of generative AI. *Computers and Education: Artificial Intelligence* 7 (2024), 100294.
- [7] Dean W. Ball and Ketan Ramakrishnan. 2025. Entity-Based Regulation in Frontier AI Governance. *Carnegie Endowment for International Peace* (2025). <https://carnegieendowment.org/research/2025/06/artificial-intelligence-regulation-united-states?lang=en> Published July 7, 2025.
- [8] Robert Birnbaum and Paul J Edelson. 1989. How colleges work: The cybernetics of academic organization and leadership.
- [9] William Blackstone. 1765. *Commentaries on the Laws of England*. Clarendon Press, Oxford. Book 4, Chapter 27 has the quote..
- [10] Rishi Bommasani, Kevin Klyman, Shayne Longpre, Sayash Kapoor, Nestor Maslej, Betty Xiong, Daniel Zhang, and Percy Liang. 2023. The Foundation Model Transparency Index. *arXiv:2310.12941 [cs.LG]* <https://arxiv.org/abs/2310.12941>
- [11] Marcus Alphons Petrus Bovens, Robert E Goodin, and Thomas Schillemans. 2014. *The Oxford handbook of public accountability*. Oxford handbooks.
- [12] Madalina Busuioc. 2021. Accountable artificial intelligence: Holding algorithms to account. *Public administration review* 81, 5 (2021), 825–836.
- [13] Cengage Group. 2024. *2024 Graduate Employability Report*. Technical Report. Cengage Group. <https://cengage.widen.net/s/bmjxxjx9mm/cg-2024-employability-survey-report>
- [14] Center for Innovative Teaching and Learning. 2024. AI Detectors: An Ethical Minefield. *NIU Center for Innovative Teaching and Learning News* (2024). <https://citl.news.niu.edu/2024/12/12/ai-detectors-an-ethical-minefield/>
- [15] Cecilia Ka Yuk Chan. 2023. A comprehensive AI policy education framework for university teaching and learning. *International journal of educational technology in higher education* 20, 1 (2023), 38.
- [16] Yucheng Chu, Hang Li, Kaiqi Yang, Harry Shomer, Yasemin Copur-Gencturk, Leonora Kaldaras, Kevin Haudek, Joseph Krajcik, Namsoo Shin, Hui Liu, and Jiliang Tang. 2025. A LLM-Powered Automatic Grading Framework with Human-Level Guidelines Optimization. In *Proceedings of the 18th International Conference on Educational Data Mining*, Caitlin Mills, Giora Alexandron, Davide Taibi, Giosuè Lo Bosco, and Luc Paquette (Eds.). International Educational Data Mining Society, Palermo, Italy, 31–41. doi:10.5281/zenodo.15870201
- [17] Simon Coghlan, Tim Miller, and Jeannie Paterson. 2021. Good proctor or “big brother”? Ethics of online exam supervision technologies. *Philosophy & Technology* 34, 4 (2021), 1581–1606.
- [18] Glenn Colby. 2025. Data Snapshot: Tenure and Contingency in US Higher Education, Fall 2023. *Academe* 111, 2 (Spring 2025). <https://www.aaup.org/academe/issues/spring-2025/data-snapshot-tenure-and-contingency-us-higher-education-fall-2023>
- [19] Common Sense Media. 2024. The Dawn of the AI Era: Teens, Parents, and the Adoption of Generative AI at Home and School. [https://www.common Sense Media.org/sites/default/files/research/report/2024-the-dawn-of-the-ai-era\\_final-release-for-web.pdf](https://www.common Sense Media.org/sites/default/files/research/report/2024-the-dawn-of-the-ai-era_final-release-for-web.pdf)
- [20] Rianne Conijn, Ad Kleingeld, Uwe Matzat, and Chris Snijders. 2022. The fear of big brother: The potential negative side-effects of proctored exams. *Journal of Computer Assisted Learning* 38, 6 (2022), 1521–1534.
- [21] Zara Contractor and Germán Reyes. 2025. Generative AI in Higher Education: Evidence from an Elite College. *arXiv:2508.00717 [econ.GN]* <https://arxiv.org/abs/2508.00717>
- [22] Cornell University Administrative AI Task Force. 2024. *Generative AI in Administration: Cornell University Task Force Report*. Technical Report. Cornell University. [https://it.cornell.edu/sites/default/files/itc-drupal10-files/Cornell%20Generative%20AI%20in%20Administration%20Task%20Force%20Report-Jan2024\\_1.pdf](https://it.cornell.edu/sites/default/files/itc-drupal10-files/Cornell%20Generative%20AI%20in%20Administration%20Task%20Force%20Report-Jan2024_1.pdf)
- [23] Ahmed M Elkhata, Khaled Elsaid, and Saeed Almeer. 2023. Evaluating the efficacy of AI content detection tools in differentiating between human and AI-generated text. *International Journal for Educational Integrity* 19, 1 (2023), 17.
- [24] EngageAI Bill of Rights Task Force. 2024. *An AI Bill of Rights for Educators*. Technical Report. EngageAI Institute.
- [25] Ethan Tuttle. 2026. LSU Students Face Mounting AI Cheating Allegations; Detection Technology Creates Campus Controversy. WAFB Channel 9 News. <https://www.wafb.com/2026/01/14/lsu-students-face-mounting-ai-cheating-allegations-detection-technology-creates-campus-controversy/>

- [26] Loredana Fattorini, Nestor Maslej, Raymond Perrault, Vanessa Parli, John Etchemendy, Yoav Shoham, and Katrina Ligett. 2024. *The Global AI Vibrancy Tool*. Technical Report. AI Index Project, Institute for Human-Centered AI, Stanford University. [https://hai.stanford.edu/assets/files/global\\_ai\\_vibrancy\\_tool\\_paper\\_november2024.pdf](https://hai.stanford.edu/assets/files/global_ai_vibrancy_tool_paper_november2024.pdf) November 2024 edition.
- [27] Rafael Ferreira Mello, Cleon Pereira Junior, Luiz Rodrigues, Filipe Dwan Pereira, Luciano Cabral, Newarney Costa, Geber Ramalho, and Dragan Gasevic. 2025. Automatic Short Answer Grading in the LLM Era: Does GPT-4 with Prompt Engineering beat Traditional Models?. In *Proceedings of the 15th International Learning Analytics and Knowledge Conference (LAK '25)*. Association for Computing Machinery, New York, NY, USA, 93–103. doi:10.1145/3706468.3706481
- [28] Colleen Flaherty. 2025. How AI Is Changing—Not ‘Killing’—College. *Inside Higher Ed* (29 Aug. 2025). <https://www.insidehighered.com/news/students/academics/2025/08/29/survey-college-students-views-ai>
- [29] University Center for Teaching and University of Pittsburgh Learning. 2025. Encouraging Academic Integrity. Online resource. <https://teaching.pitt.edu/resources/encouraging-academic-integrity/>
- [30] Josh Freeman. 2025. *Student Generative AI Survey 2025*. Policy Note 61. Higher Education Policy Institute (HEPI). <https://www.hepi.ac.uk/wp-content/uploads/2025/02/HEPI-Kortext-Student-Generative-AI-Survey-2025.pdf>
- [31] Lon Luvois Fuller. 1964. *The Morality of Law*. Yale University Press, New Haven, CT.
- [32] Maya Georgieva, Jason Webb, John Stuart, John Bell, Steven Crawford, and Beth Ritter-Guth. 2025. *AI Ethical Guidelines*. Working Group Paper. EDUCAUSE. <https://library.educause.edu/resources/2025/6/ai-ethical-guidelines>
- [33] Government Technology / Tribune News Service. 2026. Student Sues University of Michigan Over AI Misconduct Accusation. Government Technology. <https://www.govtech.com/education/higher-ed/student-sues-university-of-michigan-over-ai-misconduct-accusation>
- [34] Mohammad Hadra, Karleen Cambridge, and Mostefa Mesbah. 2026. Evaluating the accuracy and reliability of AI content detectors in academic contexts. *International Journal for Educational Integrity* 22, 1 (2026), 4.
- [35] Andrés Herrera-Poyatos, Javier Del Ser, Marcos López de Prado, Fei-Yue Wang, Enrique Herrera-Viedma, and Francisco Herrera. 2025. Responsible Artificial Intelligence Systems: A Roadmap to Society’s Trust through Trustworthy AI, Auditability, Accountability, and Governance. arXiv:2503.04739 [cs.CY] <https://arxiv.org/abs/2503.04739>
- [36] Brian T. Horowitz. 2026. An Overview of AI Governance in Education. EdTech Magazine. <https://edtechmagazine.com/higher/article/2026/02/overview-ai-governance-education-perfcon>
- [37] Panos Ipeirotis. 2025. *Fighting Fire with Fire: Scalable Personalized Oral Exams with an ElevenLabs Voice AI Agent*. <https://www.behind-the-enemy-lines.com/2025/12/fighting-fire-with-fire-scalable-oral.html>
- [38] Sheila Jasanoff et al. 2004. *States of knowledge*. Taylor & Francis Abingdon, UK.
- [39] Margot E. Kaminski and Jennifer M. Urban. 2021. The Right to Contest AI. *Columbia Law Review* 121, 7 (2021). <https://www.columbialawreview.org/content/the-right-to-contest-ai/>
- [40] Michael Katell, Meg Young, Dharma Dailey, Bernease Herman, Vivian Guetler, Aaron Tam, Corinne Bintz, Daniella Raz, and P. M. Krafft. 2020. Toward situated interventions for algorithmic equity: lessons from the field. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (Barcelona, Spain) (FAT\* '20)*. Association for Computing Machinery, New York, NY, USA, 45–55. doi:10.1145/3351095.3372874
- [41] Waleed Khattak. 2025. Algorithmic Due Process Audits: A Legal Framework for AI Accountability in Government Decisionmaking. *Available at SSRN 5362310* (2025).
- [42] Niamh Kinchin. 2024. “Voiceless”: the procedural gap in algorithmic justice. *International Journal of Law and Information Technology* 32 (2024), eaae024.
- [43] Weixin Liang, Mert Yuksekgonul, Yining Mao, Eric Wu, and James Zou. 2023. GPT detectors are biased against non-native English writers. *Patterns* 4, 7 (2023), 100779. doi:10.1016/j.patter.2023.100779
- [44] Stefanie A Lindquist and Frank C Cross. 2010. Stability, predictability and the rule of law: stare decisis as reciprocity norm. In *Conference on Legal Decision Making (Indiana)*.
- [45] Henrietta Lyons, Eduardo Velloso, and Tim Miller. 2021. Conceptualising contestability: Perspectives on contesting algorithmic decisions. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW1 (2021), 1–25.
- [46] Florence Martin, Drew Polly, Shanna Coles, and Chuang Wang. 2020. Examining higher education faculty use of current digital technologies: Importance, competence, and motivation. *International Journal of Teaching and Learning in Higher Education* 32, 1 (2020), 73–86.
- [47] Nora McDonald, Aditya Johri, Areej Ali, and Aayushi Hingle Collier. 2025. Generative artificial intelligence in higher education: Evidence from an analysis of institutional policies and guidelines. *Computers in Human Behavior: Artificial Humans* 3 (2025), 100121.
- [48] Tortoise Media. 2024. The Global Artificial Intelligence Index: Methodology Report. [https://www.tortoisemedia.com/\\_app/immutable/assets/AI-Methodology-2409.BGTLUPC-.pdf](https://www.tortoisemedia.com/_app/immutable/assets/AI-Methodology-2409.BGTLUPC-.pdf)
- [49] Jacob Metcalf, Emanuel Moss, Elizabeth Anne Watkins, Ranjit Singh, and Madeleine Clare Elish. 2021. Algorithmic impact assessments and accountability: The co-construction of impacts. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 735–746.

- [50] Benjamin Luke Moorhouse, Marie Alina Yeo, and Yuwei Wan. 2023. Generative AI tools and assessment: Guidelines of the world's top-ranking universities. *Computers and Education Open* 5 (2023), 100151.
- [51] Alexi Mostrous, Serena Cesareo, and Joe White. 2024. The Global Artificial Intelligence Index 2024. *Tortoise Media* (Sep 2024). <https://www.tortoisemedia.com/2024/09/18/the-global-artificial-intelligence-index-2024>
- [52] Matti Mäntymäki, Matti Minkkinen, Teemu Birkstedt, and Mika Viljanen. 2023. Putting AI Ethics into Practice: The Hourglass Model of Organizational AI Governance. arXiv:2206.00335 [cs.AI] <https://arxiv.org/abs/2206.00335>
- [53] National Endowment for the Humanities. 2025. Humanities Perspectives on Artificial Intelligence. <https://www.neh.gov/AI>.
- [54] National Institute of Standards and Technology. 2023. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST Special Publication. National Institute of Standards and Technology (NIST). <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf>
- [55] Pablo Fuentes Nettel, Emma Hankins, Richard Stirling, Giulia Cirri, Gonzalo Grau, Sulamaan Rahim, and Eddie Crampton. 2024. *Government AI Readiness Index 2024*. Report 7th Edition. Oxford Insights.
- [56] Kavitha Palaniappan, Elaine Yan Ting Lin, and Silke Vogel. 2024. Global regulatory frameworks for the use of artificial intelligence (AI) in the healthcare services sector. In *Healthcare*, Vol. 12. MDPI, 562.
- [57] Emmanouil Papagiannidis, Patrick Mikalef, and Kieran Conboy. 2025. Responsible artificial intelligence governance: A review and research framework. *The Journal of Strategic Information Systems* 34, 2 (2025), 101885.
- [58] Britt S. Paris, Cynthia Conti-Cook, Daniel Greene, Kyle M. L. Jones, Brian Justie, Matthew Kirschenbaum, Lisa Kresge, Emma May, Aiha Nguyen, Rebecca Reynolds, et al. 2025. Artificial Intelligence and Academic Professions. *Academe* 111, 3 (July 2025), 49–59.
- [59] Mike Perkins, Jasper Roe, Binh H Vu, Darius Postma, Don Hickerson, James McGaughran, and Huy Q Khuat. 2024. GenAI detection tools, adversarial techniques and implications for inclusivity in higher education. *arXiv preprint arXiv:2403.19148* (2024).
- [60] Jessica Priest. 2025. Texas universities deploy AI tools to review and rewrite how some courses discuss race and gender. *The Texas Tribune* (15 Dec. 2025). <https://www.texastribune.org/2025/12/15/texas-universities-ai-course-audits/>
- [61] Lee Rainie and C. Edward Watson. 2026. *The AI Challenge: How College Faculty Members Assess the Present and Future*. Technical Report. Elon University's Imagining the Digital Future Center and American Association of Colleges and Universities (AAC&U). <https://imaginingthedigitalfuture.org/wp-content/uploads/2026/01/Elon-AACU-faculty-AI-survey-full-report-1-21-26.pdf>
- [62] Marie T Reilly. 2015. Due process in public university discipline cases. *Penn St. L. Rev.* 120 (2015), 1001.
- [63] Respondus, Inc. 2026. Respondus. Respondus, Inc. <https://web.respondus.com/>
- [64] Teresa Scassa. 2022. The surveillant university: Remote proctoring, AI, and human rights. *Can. J. Comp. & Contemp. L.* 8 (2022), 271.
- [65] Andrew D Selbst. 2021. An institutional view of algorithmic impact assessments. *Harv. J.L. & Tech.* 35 (2021), 117.
- [66] Andrew D. Selbst, Danah Boyd, Sorelle A. Friedler, Suresh Venkatasubramanian, and Janet Vertesi. 2019. Fairness and Abstraction in Sociotechnical Systems. In *Proceedings of the Conference on Fairness, Accountability, and Transparency* (Atlanta, GA, USA) (FAT\* '19). Association for Computing Machinery, New York, NY, USA, 59–68. doi:10.1145/3287560.3287598
- [67] Lucy A. Suchman. 2007. *Human-Machine Reconfigurations: Plans and Situated Actions* (2nd ed.). Cambridge University Press, Cambridge, UK.
- [68] Harini Suresh, Emily Tseng, Meg Young, Mary Gray, Emma Pierson, and Karen Levy. 2024. Participation in the age of foundation models. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*. 1609–1621.
- [69] Araz Taeiagh. 2021. Governance of artificial intelligence. *Policy and society* 40, 2 (2021), 137–157.
- [70] The Hoya. 2025. Regulated AI Use, Detection at GU Leads to Growing Uncertainty. <https://thehoya.com/features/regulated-ai-use-detection-at-gu-leads-to-growing-uncertainty/>.
- [71] Turnitin. 2023. Understanding false positives within our AI writing detection capabilities. <https://www.turnitin.com/blog/understanding-false-positives-within-our-ai-writing-detection-capabilities>.
- [72] United Nations System Chief Executives Board for Coordination. 2024. *UN System White Paper on AI Governance*. Technical Report. UN System Chief Executives Board for Coordination. <https://unsceb.org/sites/default/files/2024-11/UNSystemWhitePaperAIGovernance.pdf> White Paper, November 2024.
- [73] University of Michigan Artificial Intelligence Research Committee. 2024. *Artificial Intelligence Research Committee Recommendations Report*. Technical Report. University of Michigan, Office of the Vice President for Research. <https://research.umich.edu/wp-content/uploads/2024/11/AI-Report-2024.pdf>
- [74] U.S. Census Bureau. n.d.. Census Regions and Divisions of the United States. [https://www2.census.gov/geo/pdfs/maps-data/maps/reference/us\\_regdiv.pdf](https://www2.census.gov/geo/pdfs/maps-data/maps/reference/us_regdiv.pdf).
- [75] James D. Walsh. 2025. Everyone Is Cheating Their Way Through College: ChatGPT has unraveled the entire academic project. *Intelligencer (New York Magazine)* (7 May 2025). <https://nymag.com/intelligencer/article/openai-chatgpt-ai-cheating-education-college-students-school.html>
- [76] Hui Wang, Anh Dang, Zihao Wu, and Son Mac. 2024. Generative AI in higher education: Seeing ChatGPT through universities' policies, resources, and guidelines. *Computers and Education: Artificial Intelligence* 7 (2024), 100326.
- [77] Debora Weber-Wulff, Alla Anohina-Naumeca, Sonja Bjelobaba, Tomás Foltýnek, Jean Guerrero-Dib, Olumide Popoola, Petr Šigut, and Lorna Waddington. 2023. Testing of detection tools for AI-generated text. *International Journal for Educational Integrity* 19, 1 (2023),

1–39.

- [78] Karl E Weick. 2012. Educational organizations as loosely coupled systems. In *The roots of logistics*. Springer, 389–410.
- [79] Langdon Winner. 1980. Do Artifacts Have Politics? *Daedalus* (1980), 121–136.
- [80] Chuhao Wu, He Zhang, and John M. Carroll. 2024. AI Governance in Higher Education: Case Studies of Guidance at Big Ten Universities. arXiv:2409.02017 [cs.HC] <https://arxiv.org/abs/2409.02017>
- [81] Junchao Wu, Shu Yang, Runzhe Zhan, Yulin Yuan, Lidia Sam Chao, and Derek Fai Wong. 2025. A survey on llm-generated text detection: Necessity, methods, and future directions. *Computational Linguistics* 51, 1 (2025), 275–338.
- [82] Meg Young, Michael Katell, and PM Krafft. 2022. Confronting power and corporate capture at the FAccT Conference. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*. 1375–1386.
- [83] Jiawei Zhang and Boris Paal. 2025. A Three-Layered Framework: An AI Governance Guide for Global Policymakers. (2025). <https://papers.ssrn.com/abstract=5241351>

## A Positionality Considerations

Our interdisciplinary team spans computer science, engineering, humanities, and information science. This collaboration has been central to shaping this work, bringing in both the computational and humanistic ways of knowing.

## B Prior Work on Institutional AI Governance

Prior efforts to examine how universities respond to AI take two forms. The first consists of *empirical audits*: studies that systematically analyze institutional governance practices across a defined sample of universities, shown in Table 2. The second consists of *community-maintained collections*: curated inventories of AI policies, syllabus statements, and guidelines aggregated by universities or individuals, shown in Table 3. These collections serve as valuable reference points but do not apply a systematic audit framework, annotation protocol, or comparative scoring. ACAI-US79 bridges both categories — contributing an audited dataset, a reproducible index, and public infrastructure for ongoing comparison.

## C AI Detection Tools

Independent evaluations of AI detection tools consistently find them unreliable. As summarized in Table 4, accuracy across tools rarely exceeds 80%, degrades substantially against newer language models, and drops further still when students employ adversarial techniques like paraphrasing or adding spelling errors [23, 59, 77]. Even the most widely adopted commercial tools show significant limitations: Hadra et al. [34] evaluated Turnitin and Originality across 192 texts — including authentic writing by English as a Foreign Language (EFL) students produced before the generative AI storm — and found overall accuracy of just 0.61 and 0.69, respectively, with both tools performing especially poorly on hybrid human-AI texts.

Beyond overall accuracy, a growing body of evidence documents systematic demographic disparities in false positive rates. Liang et al. [43] found that seven widely used detectors misclassified over 47% of essays by non-native English speakers as AI-generated, while maintaining less than 12% misclassification rates for native English speakers. The mechanism is perplexity-based: limited lexical diversity and syntactic complexity — common features of second-language writing — overlap with the statistical signatures these tools use to flag AI-generated text. Hadra et al. [34] conclude that “their limitations make them unsuitable as the sole basis for decisions regarding academic misconduct,” emphasizing that “for [EFL]-focused institutions in particular, over-reliance on machine-based identification risks misclassifying legitimate student work” and calling for “continued development of equitable detection methods that account for linguistic diversity.” These disparities extend beyond language background: Black teens are approximately twice as likely as their peers to report that teachers flagged their schoolwork as AI-generated when it was not [19].

Despite this evidence, universities continue to embed these tools into their governance infrastructure with less-than-adequate acknowledgment of their limitations. As shown in Table 6, Nova Southeastern University

|                                     | Wu et al. [80]       | Moorhouse et al. [50]                                                        | Wang et al. [76]                             | McDonald et al. [47]                  | ACAI (Ours)                                                                     |
|-------------------------------------|----------------------|------------------------------------------------------------------------------|----------------------------------------------|---------------------------------------|---------------------------------------------------------------------------------|
| <b>Sample</b>                       | 14 R1 (U.S.)         | Top 50 (global)                                                              | Top 100 (U.S.)                               | 116 R1 (U.S.)                         | 79 U.S. universities                                                            |
| <b>Institutional Diversity</b>      | Major R1 only        | Elite only                                                                   | Elite only                                   | R1 only                               | R1, R2, Teaching/ Liberal Arts                                                  |
| <b>Analytic Frame</b>               | Pedagogical guidance | Assessment guidelines                                                        | Policy & resource presence                   | Policy & guideline content            | Procedural accountability                                                       |
| <b>Governance Domains</b>           | AI use guidance      | Academic Integrity, Advice on Assessment Design, Communication with Students | Policy, resources, re-detection tool mention | Policy, guidelines, faculty resources | Policy Clarity, Faculty Support, Feedback Mechanisms, Detection Tool Safeguards |
| <b>Feedback &amp; Participation</b> | ✗                    | ✗                                                                            | ✗                                            | ✗                                     | ✓                                                                               |
| <b>Detection Tool Governance</b>    | Mention only         | ✓                                                                            | ✓                                            | ✓                                     | Procedural safeguards                                                           |
| <b>Scoring / Index</b>              | ✗                    | ✗                                                                            | ✗                                            | Spectrum-based placement              | ACAI (0–100)                                                                    |
| <b>Annotation</b>                   | Authors              | Authors                                                                      | Authors                                      | Authors                               | Paid annotators (Prolific) + expert validation (authors)                        |
| <b>Public Dataset</b>               | ✗                    | ✗                                                                            | ✗                                            | ✗                                     | ✓                                                                               |
| <b>Public Website</b>               | ✗                    | ✗                                                                            | ✗                                            | ✗                                     | ✓                                                                               |

Table 2. **Comparison of Prior Audits of AI Governance in Higher Education.** ACAI extends prior work by auditing a structurally diverse institutional sample and shifting the analytic frame from pedagogical guidance to procedural accountability.

specifically directs faculty to ZeroGPT, describing it as a tool where “[y]ou can easily detect whether text is human-written or AI/GPT generated” and noting that the “company claims 99% and higher accuracy rate.” When institutions surface vendor accuracy claims without reference to independent evaluations or known demographic disparities, they risk translating tool-level unreliability into policy-level harm for linguistically diverse and racially marginalized students.

Prior work has documented the unreliability of AI detection tools and called for caution in their deployment in educational settings [34, 43, 77]. Our audit of 79 university AI policies extends this line of work by examining whether these calls have been reflected in institutional practice. The evidence from ACAI-US79 suggests they largely have not. This shifts the conversation from whether detection tools are flawed — a point the literature has established — to how institutions are governing their adoption, and what safeguards are absent.

| University or Organization            | Resource                                      | Comments                                                                                                                                                                                                                                                                                                                                                                   | Link |
|---------------------------------------|-----------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|
| University of La Verne                | List of institutions with AI guidelines       | <i>"We found that most universities that we surveyed have some kind of statement or set of guidelines for genAI in the classroom. This list is far from exhaustive."</i>                                                                                                                                                                                                   | Link |
| Western University of Health Sciences | University Policies on Generative AI          | <i>"Collection of university policies and websites. Questions? Contact CETL@westernu.edu."</i>                                                                                                                                                                                                                                                                             | Link |
| Northeastern University               | A moderated list of AI syllabus statements    | <i>"If you would like to submit your course guidelines/policy or revise your submission, please submit it in this form."</i>                                                                                                                                                                                                                                               | Link |
| Northeastern University               | A moderated list of AI institutional policies | <i>"This document is maintained by Lance Eaton. You are welcome to share it with other individuals, groups, and organizations. To view the policies, please select the "Policies" tab in this spreadsheet. If you would like to submit your policy, please complete this form (https://bit.ly/AI-Institutional-Policies) and it will show up here within 24-48 hours."</i> | Link |
| Gradpilot                             | The State of AI in College Admissions         | <i>"Navigate AI usage rules across 150+ American universities"</i>                                                                                                                                                                                                                                                                                                         | Link |

Table 3. **Community-maintained collections of institutional AI policies and guidelines.** These resources aggregate policy documents, syllabus statements, and administrative guidance, but do not apply a systematic audit framework, annotation protocol, or comparative scoring. ACAI complements these collections by providing a structured, reproducible audit with paid annotation and a public index.

Concerns about plagiarism and unauthorized use of generative AI tools are often treated as isolated classroom management issues. In reality, they reveal deeper failures of governance capacity – specifically, the lack of clear institutional policies, interpretive frameworks, and governance mechanisms capable of addressing AI use in equitable, accountable, and pedagogically meaningful ways. Universities have frequently responded to the rise of generative tools with restrictive policies or punitive enforcement regimes, but these approaches are often built on vague or inconsistent definitions of permissible use. The result is widespread uncertainty among both students and faculty [75], inconsistent application of standards, and in some cases, false accusations of academic misconduct.

While it is true that many faculty are concerned with academic dishonesty with the rise of generative AI, it is also the case that for at least some faculty, the concern goes beyond the (dis)honesty question and is rooted in a worry that students will fail to learn; as the AAUP Report puts it:

*"The distinction between honesty and failure to learn is critical because it highlights one of the core goals of higher education: to develop a well-informed and thoughtful citizenry. This finding suggests that there is a need for higher education to refocus on the relational aspects of education and learning, as opposed to punitive measures..."* – Paris et al. [58]

Such outcomes erode trust, undermine institutional legitimacy, and highlight the urgent need for governance structures that are transparent, participatory, and aligned with civic values.

Table 4. Selected evaluations of AI detection tool accuracy and documented demographic disparities in false positive rates.

| Detection Tool Accuracy |                                 |                                                                                                                                            |                                                                                                                              |
|-------------------------|---------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|
| Study                   | Scope                           | Key Finding                                                                                                                                | Details                                                                                                                      |
| Weber-Wulff et al. [77] | 14 tools tested                 | All tools <80% accuracy                                                                                                                    | Turnitin and then Compilatio scored highest. Adversarial techniques like machine translation, paraphrasing reduced accuracy. |
| Perkins et al. [59]     | 6 tools tested                  | All tools <75% accuracy                                                                                                                    | Adversarial techniques like adding spelling errors, increasing complexity reduced accuracy.                                  |
| Elkhatat et al. [23]    | 5 tools tested                  | More accurate for GPT-3.5 than GPT-4                                                                                                       | Authors discuss that models change quickly, and detectors struggle to keep up.                                               |
| Hadra et al. [34]       | 2 tools (Turnitin, Originality) | Both tools <70% accuracy                                                                                                                   | Dataset of 192 texts which include EFL student writing.                                                                      |
| Liang et al. [43]       | 7 tools tested                  | All tools misclassify (at >47%) <b>non-native english speaker</b> essays as AI-generated, in contrast to < 12% for native english speakers | Used adversarial prompting ( <i>Enhance the word choices to sound more like that of a native speaker.</i> ) to probe.        |

The stakes of false accusations extend far beyond individual classroom incidents. Let us consider Blackstone's Ratio [9]:

*"It is better that ten guilty persons escape than one innocent suffer."*

This principle remains instructive in the context of AI governance. **Institutions that tolerate high rates of false positives in AI detection – for example, by punishing students based on unreliable systems – risk delegitimizing their authority and weakening the educational contract itself.** Even a 1% false positive rate can result in thousands of wrongful accusations across large institutions, eroding trust in both faculty and administrative oversight, as illustrated in Figure 4. Protecting the innocent is foundational to sustaining the *legitimacy* of the educational institution as a whole.

Moreover, reliance on flawed detection technologies [81] introduces new forms of procedural injustice. Many detection systems disproportionately misclassify the work of multilingual students or those who write in nonstandard styles, compounding existing inequities [14, 43]. At the same time, the usefulness in certain contexts and ubiquity of generative AI virtually guarantee its continued presence in certain academic and professional contexts [6]. Attempts to suppress use via detection technologies are both impractical and counterproductive. Instead, educators should focus on equipping students with the critical skills needed to evaluate, contextualize, and responsibly integrate (or not integrate) AI-generated content into their work.

These challenges highlight that preventing misuse is not simply a question of enforcement but one of governance design. Effective institutional responses will require clear and consistently applied definitions of acceptable AI use, transparent policies with built-in safeguards such as appeal and review mechanisms, and curriculum that teaches students how to engage with generative AI ethically and productively. By shifting from a reactive,



Fig. 4. **False Positive Rate of 1%:** As an example, a false positive rate of only 1% [29, 71] indicates that 1 in every 100 cases will be flagged for cheating incorrectly. It falls on the instructor to evaluate use – faculty who are ill-equipped to assess state-of-the-art AI systems. At scale in a university setting, for every 100,000 submissions, 1000 false accusations could occur, jeopardizing student careers and significantly de-legitimizing educational institutions.

punitive stance to a proactive governance model, universities can both uphold academic integrity and prepare students for meaningful participation in society increasingly pervaded by AI.

#### D Public Website

To support transparency, inspection, and reuse of the ACAI-US79 audit, we provide a public, interactive website at <https://acai-us79.org>. The site is designed to make both the audit results and the audit instrument accessible beyond the paper, enabling readers to explore institutional AI governance capacity and to trace aggregate scores back to the underlying publicly available materials.

As shown in Figure 5, the primary interface is an interactive map-based visualization of the 79 U.S. universities included in the audit. Each institution is represented as a clickable marker, and a synchronized sidebar lists universities ranked by ACAI score. Users can filter institutions by research activity, institutional type, geographic region, and student size, and can select either a map marker or a list entry to reveal institution-specific details. For each university, the interface displays its ACAI score alongside direct links to the institutional policies, teaching resources, governance committees, and academic integrity materials reviewed during the audit. This design explicitly supports traceability, allowing users to inspect how scores are grounded in publicly legible governance artifacts rather than treating the index as an opaque ranking.

The website also includes a self-scoring interface for institutions not included in the ACAI-US79 dataset. This tool implements the same annotation schema and scoring criteria used in the audit, enabling users to assess their own institution using publicly available materials under comparable constraints. Self-scoring results are returned solely for informational and reflective purposes and are not incorporated into the released dataset.

The application is implemented in React with TypeScript, using `react-simple-maps` for geographic visualization and `Material-UI` for interface components. Institutional metadata—including location coordinates, research classification, institutional type, and associated governance resources—is stored in a structured JSON format




Fig. 5. **Interactive map of ACAI-US79** at <https://acai-us79.org/>, visualizing the 79 U.S. universities and describing their publicly articulated governance capacity. Institutions are shown as clickable markers and ranked by ACAI score, with filters enabling comparison across research activity, institutional type, region, and size. Selecting an institution reveals its score and links to the publicly available policies, guidance, and governance materials reviewed in the audit, supporting traceability and independent inspection of how governance capacity is publicly articulated.

derived from the original audit datasets. The interface supports interactive filtering, zooming, and state-level annotations to facilitate exploration and comparison across institutions.

Overall, the website is designed to emphasize interpretability and accountability. ACAI scores reflect the public legibility of institutional governance artifacts at the time of review and should not be interpreted as measures of internal practice or intent. By making both the audit results and the audit instrument publicly accessible, the site supports replication, critique, and institutional self-reflection, and positions ACAI as an auditable, contestable infrastructure rather than a static evaluation.

## E Extended Related Work

### E.1 Institutional & Socio-Technical Systems Theory, and Algorithmic Accountability

As Winner argued, “*What matters is not technology itself, but the social or economic system in which it is embedded*” [79]. This position, described as social determination theory, offers a corrective to naïve technological determinism: the assumption that technology evolves according to its own internal logic and subsequently shapes society in a one-directional way. Technological and institutional capacities do not evolve independently, instead co-producing

and mutually constructing each other [38]. This aligns with the argument of Selbst et al. [66], who warn that abstracting away from social context obscures the structural forces that shape fairness itself.

Algorithmic accountability is usually determined via algorithmic impact assessments, an auditing mechanism to judge how an algorithm is causing harm. Metcalf et al. [49] highlights the distinctiveness of algorithmic systems, noting that there is a heightened risk with misunderstanding the inner workings of algorithmic systems, and therefore being unable to legislate their use or development effectively. Selbst [65] also note this gap between the inner workings of a computational system and effective governance, demonstrating cases of “algorithmic harm where existing liability regimes fail to hold the creators of the harm to account, specifically because of a lack of knowledge about the development process.” Notably, Ananny and Crawford [4] criticize the transparency ideal in algorithmic assessments, noting that “transparency alone cannot create accountable systems.”

## E.2 AI Governance Models

Recently, there have been calls for attention to AI governance [56, 57, 69]. Prior AI governance models formulate the relationship between governance and computation as unidirectional capacity flows, as shown in Table 5. Across this literature, governance and computation are treated as sequential processes, where one directs, regulates, or reacts to the other. These frameworks have advanced the field’s understanding of risk and ethics, but they remain largely reactive and linear, conceptualizing governance as either a *top-down constraint* or a *downstream response* to technical innovation.

Existing models recognize the importance of institutional capacity: The UN System Survey of Institutional Models [72] emphasizes “capacity-building” but defines it primarily as the technical training of scientists and regulators; The Responsible AI Systems Roadmap [35] focuses on the role of scientists in shaping policy. While models recognize the importance of institutional development, they largely equate capacity with technical skill or regulatory compliance rather than with civic or educational infrastructure.

**Governance Capacity → Computational Capacity.** The Hourglass Model of Organizational AI Governance [52] represents an ethics-first pipeline in which governance capacity flows toward computation. Ethical principles and organizational oversight mechanisms are translated into engineering practice. While there are feedback mechanisms for communication, the model largely presumes institutions can directly steer technical behavior. A similar directional logic appears in the Entity-Based Regulation Framework [7], which centers regulatory oversight of large corporate actors, confined to the level of enforcement and transparency rather than broader institutional design. Both frameworks view governance as an initiating force, with the Hourglass model supporting a limited feedback mechanism.

**Computational Capacity → Governance Capacity.** Other frameworks reverse this flow, positioning technical development as the driver of governance. The NIST AI Risk Management Framework [54] sequences accountability as mapping, measuring, and managing risk and emphasizes flexibility, voluntarism, and scalability across sectors:

*“The Framework is intended to be voluntary, rights-preserving, non-sector-specific, and use-case agnostic, providing flexibility to organizations of all sizes and in all sectors and throughout society to implement the approaches in the Framework.”*

The Three-Layered Framework [83] follows a similar pipeline but extends it globally, envisioning governance as a set of layered responses to computational markets – drawing from the toolbox categories of market-invigorating strategies, value-directed rules, and procedural controls. The UN System Survey of Institutional Models [72] emphasizes “capacity-building” but defines it primarily as the technical training of scientists and regulators. The Responsible AI Systems Roadmap [35] focuses on the role of scientists in shaping policy. These

| Model                                                    | Primary Focus                                                                                                                                                                                                             | Capacity Flow              | Scope     |
|----------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------|-----------|
| The Hourglass Model of Organizational AI Governance [52] | Focused on AI ethics, risk mitigation; organized into environmental, organizational, and AI system layers with some feedback mechanisms (i.e. computational → governance).                                                | Governance → Computational | EU/Global |
| NIST AI Risk Management Framework [54]                   | Focused on identifying AI risks, which governance Maps, Measures, and Manages.                                                                                                                                            | Computational → Governance | USA       |
| Entity-Based Regulation Framework [7]                    | Focused on transparency and regulating “ <i>the large business entities developing the most powerful AI models and systems</i> ” with emphasis on preemptive risk regulation.                                             | Governance → Computational | USA       |
| Three-Layered Framework [83]                             | Focused on risk, aims to fix market failures using a toolbox of regulatory tools from the three layers: market-invigorating strategies, value-directed rules, and procedural controls.                                    | Computational → Governance | Global    |
| UN System Survey of Institutional Models [72]            | Focused on ethics and risk; highlights “ <i>[national] capacity-building [that] can support AI development that is grounded in fairness, gender equality, reliability, safety, interpretability and accountability.</i> ” | Computational → Governance | UN/Global |
| Responsible AI Systems Roadmap [35]                      | Focused on risk, highly dependent on a committee of scientists to shape policy.                                                                                                                                           | Computational → Governance | UN/Global |
| AI Ecological Education Policy Framework [15]            | Focused on education; organized into 3 dimensions of educational support: pedagogical, ethical, and operational.                                                                                                          | Computational → Governance | Hong Kong |

Table 5. Existing governance frameworks recognize that there is both *computational capacity* and *governance capacity*.

models recognize the importance of institutional development yet largely equate capacity with technical skill or regulatory compliance rather than with civic or educational infrastructure.

### E.3 Institutional Capacity

Institutional Capacity extends beyond Governance Capacity to include the civic, educational, industrial, and governmental infrastructures that make computational work socially legitimate and accountable. It encompasses the systems and organizations that translate technical advancement into legal systems, economic sectors, government, defense, and culture, and therefore requires a wide and interdisciplinary range of expertise:




Fig. 6. **Overview of Iterative Link Categorization Process:** See Figure 7 for more details on Phases II and III. Green highlights show refinements.

*“Questions about the impact of AI on American society and culture are fundamentally rooted in such humanities fields as ethics, law, history, philosophy, anthropology, sociology, media studies, and cultural studies.”* –National Endowment for the Humanities [53]

While computational capacity has been extensively benchmarked, its rapid expansion has outpaced the institutional systems needed for accountability.

#### E.4 The Translation Workforce

The translation workforce is professionals who bridge technical and institutional domains. Algorithmic action is inherently situated – systems can have unintended outcomes when deployed in real social and organizational contexts [40, 67]. Effective governance therefore requires people and processes capable of translating between computational reasoning and institutional judgment.

The translation workforce goes beyond AI literacy (understanding how AI operates) or ethics (identifying harms). It encompasses the practical work of institutional integration: drafting regulations, designing oversight mechanisms, adjudicating disputes, negotiating standards, and explaining algorithmic decisions to diverse stakeholders, and/or exercising judgment about when AI implementation is not appropriate. Engineers, policymakers, educators, artists, cultural critics, and legal professionals all participate in this interpretive process.

#### F ACAI-US79 Iterative Link Categorization Process

This section describes the iterative link categorization process used to construct ACAI-US79: Figure 6 gives an overview of the process, and Figure 7 shows the detailed annotator instructions for Phases II and III. Because institutional AI governance is unevenly distributed and inconsistently labeled across universities, we employ a recall-oriented, multi-phase procedure that surfaces publicly legible, institutionally authoritative materials while enforcing a clear boundary between governance capacity and general AI research or outreach.

In **Phase I**, as shown in Figure 6, we use an LLM to generate candidate links for each of seven target categories, which we then manually screen to remove errors, hallucinations, and off-scope pages. During manual screening, we remove dead links, incorrect links (e.g., pages from the wrong university), AI-hallucinated links that do not exist, news links, and announcement pages. This screening step ensures that the resulting corpus contains only live, correctly-attributed, on-scope institutional pages suitable for downstream analysis.

The seven categories (T1–T7) are neither mutually exclusive nor collectively exhaustive. In practice, AI governance materials are unevenly distributed across institutional units or consolidated into a small number of

centralized resources. As a result, a single link may be relevant to multiple categories, while other categories may contain no links for a given institution. Such empty cells reflect genuine variation in how AI governance is organized and communicated, not missing data or annotation error. Category membership should therefore be interpreted as evidence of where—and how legibly—governance functions are articulated in public-facing materials.

**Phases I and II** are deliberately recall-oriented and surface a wide range of AI-related materials, including research and outreach content. **Phase III** enforces the conceptual boundary of the audit by removing links that are AI-related but governance-irrelevant, such as research labs, grants, or faculty-led initiatives that do not articulate institutional authority or procedural expectations. This consolidation step prevents research-intensive institutions from appearing more “governed” simply due to higher volumes of AI-related content and ensures that institutional capacity is not conflated with research productivity. The resulting link set retains only publicly legible, institutionally authoritative governance artifacts, such as standing policies, centrally maintained guidance, and formal committees. By narrowing annotator attention to governance-relevant materials, Phase III also improves annotation consistency and strengthens the construct validity of the audit.

We do not include the following types of links, because we found in Phases I and II of our iterative link categorization process that these types of links generally weren’t relevant to our annotation criteria (\*) or were separate from our focus in this work (+):

- **Legal / General Counsel\***: Especially for risk, compliance, data use, copyright, FERPA, and contracts involving AI tools. We also think that students have access to a variety of AI tools separate from those visible by the university.
- **Research Office / Office of Sponsored Programs<sup>+</sup>**: Research oversight is mediated through specialized mechanisms (e.g., AI use in grants, data management plans, human subjects, and responsible research conduct) separate from the present audit, which focuses on governance as enacted through core university functions shaping teaching, assessment, and student experience.
- **Data Governance / Privacy Office\***: Sometimes separate from IT; increasingly relevant for AI training data and student data use.
- **Accessibility / Disability Services<sup>+</sup>**: AI accommodations, assistive tech, and equity considerations. While critically important for AI equity, guidance in this area is often individualized/case-specific.
- **Human Resources<sup>+</sup>**: Staff and faculty use of AI for hiring, evaluation, or administrative work.
- **Admissions<sup>+</sup>**: Policies on AI-assisted application materials are often separate from academic integrity rules.
- **Graduate School<sup>+</sup>**: Graduate-specific guidance often differs from undergraduate rules.
- **Department or College-Level Pages<sup>+</sup>**: Many institutions defer AI guidance to colleges (e.g., Engineering, Business, Law) or even individual departments.

This dataset reflects publicly available, institutionally maintained web resources and therefore has a few key limitations. First, universities differ substantially in how AI-related guidance is organized, labeled, and distributed across administrative units, which may lead to uneven coverage across institutions. We aimed to yield equitable coverage with our iterative link categorization process (Figure 6) as to correctly capture the heterogeneous ways in which different universities – varying in region, research activity, student population size, and more as discussed in §3.1 – share and present AI policy. Second, the dataset captures only formal, publicly visible governance artifacts and does not reflect informal practices, internal guidance, or unpublished decision-making processes that may substantially shape how AI is used and regulated within institutions. As a result, the dataset should be interpreted as a representation of officially articulated AI governance rather than a complete account of institutional practice.

**Phase II: LLM Cleanup**

For **each university and each category (T1–T7)**, apply a conservative cleanup pass, removing links **only when they are clearly out of scope**. When determining relevance would require nontrivial investigation beyond initial screening, we conservatively retain the link and defer judgment to the annotation process, where deeper inspection is already required. This typically includes links that:

- Are non-authoritative (i.e., not published by the university)
- Are dead, redirecting, or hallucinated
- Are narrowly subject, department, or graduate school specific rather than university-level
- Are strictly research-focused
- Consist solely of news, announcements, or event listings
- Provide application instructions for prospective students
- Are intended for university communications or marketing staff
- Require institutional login for access

**Additionally, when verification is straightforward:**

- Remove malformed entries (e.g., markdown artifacts instead of valid links)
- Remove excessive sublinks pointing to the same underlying page
- De-duplicate or re-categorize links

**Phase III: Link Identification**

For **each university and each category (T1–T6)**, if a clear main link has not already been identified:


- (1) **Identify the Authoritative Institutional Surface and Any Clearly AI-Related Subpages** (Skip to Step #2 for T1)
  - (a) Perform a Google search for “University” + “Category”
  - (b) Navigate to the primary institutional landing page(s)
    - (i) Review standard navigation paths (e.g., scrolling the page, examining menu bars)
    - (ii) Include any subpages that are clearly AI-related
    - (iii) If additional clearly AI-related subpages are encountered during exploration, include them
- (2) **Targeted Confirmation Search** (If no AI-related content is visible in Step #1; except T3 and T6, which explicitly reference “AI”)
  - (a) Perform a Google search for “University” + “Category” + “AI”
  - (b) Include any institutionally authoritative pages that are clearly relevant
- (3) If no suitable links are identified, **record “None” for that category**.
- (4) Mark the university as complete once all categories are reviewed.

Fig. 7. Annotation Instructions for Phases II and III of Iterative Link Categorization Process.

### F.1 Phase II vs. III Coverage Comparison

In this section, we provide details on the links retrieved in Phase II vs. III of our iterative link categorization process, characterizing how human annotators empowered by Google Search in Phase III processed the LLM output results from Phase II. Specifically, for each institution-category pair, we compare the LLM-retrieved link set  $L$  with the human-curated set  $H$  using set-based metrics. We compute precision ( $|H \cap L|/|L|$ ) and recall ( $|H \cap L|/|H|$ ), along with the number of *added* links ( $|L \setminus H|$ ) and *deleted* links ( $|H \setminus L|$ ). We interpret recall as a measure of *coverage*, indicating whether the LLM retrieves policy evidence aligned with human judgment.

*How good is LLM coverage of institutional AI policy?* In Figure 8a, we show recall (human coverage) against precision (LLM correctness) for all institution-category pairs. The distribution shows that LLM retrieval is strongly recall-oriented: many cases achieve high recall but only moderate precision, indicating that while the LLM often retrieves at least one relevant policy link, it frequently includes additional links that are later removed by human



(a) Precision–recall characteristics of LLM-based AI policy link retrieval (Phase II) compared to human-curated links (Phase III). Each point represents one institution–category pair.

(b) LLM coverage of institutional AI policies under three definitions: raw coverage, aligned coverage, and missed coverage.

Fig. 8. Comparison of LLM-based and human-curated AI policy link retrieval.

annotators. Instances of low recall correspond to true coverage failures, where the LLM fails to retrieve any human-recognized policy evidence. To further quantify this behavior at the category level, we summarize LLM coverage using three increasingly strict definitions: *raw coverage* (the presence of any LLM-retrieved link), *aligned coverage* (overlap with human-curated links), and *missed coverage* (human-curated links not retrieved by the LLM). The results are shown in Figure 8b. While raw coverage is consistently high across categories, aligned coverage drops substantially, indicating that apparent LLM coverage overstates effective policy coverage. This discrepancy is most pronounced in categories involving decentralized or evolving institutional resources, such as libraries and AI initiatives, where authoritative policy evidence is more difficult to identify automatically. Taken together, these results suggest that LLMs perform well as a first-stage retrieval mechanism, successfully surfacing candidate AI policy evidence for most institutions and categories. However, the observed gap between apparent and aligned coverage underscores the continued necessity of human supervision to filter non-authoritative links, recover missed policy documents, and ensure that final policy representations accurately reflect institutional governance.

## F.2 Example Links

In Table 12, we provide examples of the types of links retrieved in T1-T7.

## F.3 Prolific Privacy Policy

We provide the Privacy Policy in Figure 9.

## G ACAI Domains in Practice

Table 6 provides illustrative examples of stronger and weaker institutional responses for each ACAI governance domain: **A. POLICY CLARITY**, **B. FACULTY SUPPORT**, **C. FEEDBACK MECHANISMS**, and **D. DETECTION TOOLS**. Each

Table 6. Illustrative examples of institutional responses across governance domains, drawn from ACAI-US79.

| Domain                 | Stronger Practice                                                                                                                                                                                                                                                                                                                                                                                                                                                                | Weaker Practice                                                                                                                                                                                                                                                                                                          |
|------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| A. POLICY CLARITY      | <p><i>“What is generative AI? You have probably seen or heard of ChatGPT and Claude... One of the key features of generative AI is its ability to identify patterns and structures from input data it receives, and then apply that input to automatically generate content through the use of various algorithms and models.”</i></p> <p>– Carleton College (A1, Score: 0.83)</p>                                                                                               | <p><i>“Scite and Semantic Scholar are two tools that utilize artificial intelligence to show how citations are used in context...”</i></p> <p>– Stevens Institute of Technology (A2, Score: 0.17)</p>                                                                                                                    |
| B. FACULTY SUPPORT     | <p><i>“Step 2: Choose and Define AI Usage levels. Below is starter language to help you define and describe your stance on AI use in your course. You can include and adapt it for use in your course syllabus...”</i></p> <p>– University of North Carolina at Chapel Hill (B4, Score: 0.83)</p>                                                                                                                                                                                | <p><i>“Grinnell’s policy on Honesty in Academic Work requires students to: cite all language that is quoted or paraphrased from a large language model (LLM) or other generative AI service... When in doubt, ask your instructor for guidance.”</i></p> <p>– Grinnell College (B2, Score: 0.33)</p>                     |
| C. FEEDBACK MECHANISMS | <p><i>“PSU’s Office of Academic Affairs formed a Generative AI Task Force in the Summer of 2023 that transitioned to a Working Group in the Spring of 2024. The cross-functional working group is composed of representatives from across the university, including faculty, a representative of the Faculty Senate Educational Policy Committee, the offices of Academic Innovation and Information Technology...”</i></p> <p>– Portland State University (C1, Score: 1.00)</p> | <p><i>“SMU is curating resources for faculty to explore and share about AI...”</i></p> <p>– Southern Methodist University (C1, Score: 0.33)</p>                                                                                                                                                                          |
| D. DETECTION TOOLS     | <p><i>“AI Detection Tools: Be aware that AI detection tools carry risks of misidentification and have not been widely proven to detect AI use. Careful consideration should be given when deciding if the use of these tools is appropriate for assessing student work.”</i></p> <p>– Saint Louis University (D1, Score: 0.83)</p>                                                                                                                                               | <p><i>“ZeroGPT. You can easily detect whether text is human-written or AI/GPT generated. This AI text detector displays the percentage of AI/GPT plagiarized text for an in-depth analysis of content. The company claims 99% and higher accuracy rate.”</i></p> <p>– Nova Southeastern University (D1, Score: 0.33)</p> |

example is drawn from ACAI-US79 and labeled with the specific ACAI item and average score from the  $k = 3$  annotators. These examples are intended to concretize the scoring criteria and help readers contextualize the types of policy documents evaluated in our audit.

## H ACAI Rank Changes Under Alternate Weighting Schemes

To assess whether ACAI rankings are artifacts of normative weighting choices, we evaluated four weighting schemes: an indicator-weighted baseline ( $A = 2, B = 4, C = 3, D = 2$ ), equal ( $A = 1, B = 1, C = 1, D = 1$ ), policy-heavy ( $A = 1, B = 1, C = 2, D = 2$ ), and teaching-heavy ( $A = 1, B = 2, C = 1, D = 1$ ), using percentile ranks. We show the results in Tables 7 and 8. We also assess robustness to individual annotators via a leave-one-out analysis on the equal-weighted score (Table 9).

Table 7. Pairwise Rank Correlations Across Schemes

| Scheme 1     | Scheme 2       | Spearman | Pearson |
|--------------|----------------|----------|---------|
| Baseline     | Equal          | 0.98     | 0.98    |
| Baseline     | Policy-heavy   | 0.94     | 0.94    |
| Baseline     | Teaching-heavy | 0.99     | 0.99    |
| Equal        | Policy-heavy   | 0.97     | 0.97    |
| Equal        | Teaching-heavy | 0.98     | 0.98    |
| Policy-heavy | Teaching-heavy | 0.93     | 0.93    |

Table 8. Maximum Absolute Rank Change Statistics

|              | Baseline | Equal | Policy-heavy | Teaching-heavy |
|--------------|----------|-------|--------------|----------------|
| Mean         | 16.32    | 17.32 | 16.95        | 16.51          |
| Median (50%) | 15.00    | 17.50 | 15.00        | 13.50          |
| 75%          | 22.50    | 23.75 | 22.50        | 22.25          |
| Max          | 46.00    | 46.00 | 43.50        | 45.50          |

Table 9. Leave-One-Annotator-Out Sensitivity for Indicator-Weighted ACAI Scores

| Dropped | Pearson $r$ | Spearman $\rho$ | Mean $ \Delta $ | Max $ \Delta $ |
|---------|-------------|-----------------|-----------------|----------------|
| $-k_1$  | 0.862       | 0.871           | 5.95            | 22.73          |
| $-k_2$  | 0.832       | 0.774           | 6.26            | 21.21          |
| $-k_3$  | 0.880       | 0.878           | 5.42            | 19.70          |

## I Details on Comparison to CSRankings

To contextualize institutional AI governance capacity relative to AI research activity, we constructed a reference list of research-active universities using CSRankings, shown in Figure 11, a widely used, publicly available ranking of computer science research output. CSRankings aggregates publication counts across major computer science venues and allows filtering by research area and time period.

We first navigated to [csranks.org](https://csranks.org) and restricted the ranking to AI-relevant research areas only, enabling the categories of Artificial Intelligence, Computer Vision, Machine Learning, Natural Language Processing, and The

Web & Information Retrieval, while disabling all other areas. This filtering step was intended to approximate institutional engagement with AI-related research rather than overall computer science output.

We then restricted the publication window to 2022–2025, corresponding to the period following the public release of large-scale generative AI systems (i.e., ChatGPT) and the following rapid expansion of AI use in educational contexts. This temporal filter was chosen to reflect contemporary AI research activity during the period in which universities began articulating institutional responses to generative AI.

The resulting ranked list of universities was used as a reference set for identifying institutions with substantial recent AI research activity. Importantly, this list was not treated as a measure of governance quality or institutional responsibility. Rather, it served as a comparative baseline to examine whether AI research intensity correlates with publicly articulated AI governance capacity.

This procedure relies exclusively on publicly accessible filters and settings within CSRankings and is fully reproducible.

| (a) Exclude unranked CSRankings <sub>AI</sub> ; results also visualized in Figure 10 |          |         |          |                |                            |                             |
|--------------------------------------------------------------------------------------|----------|---------|----------|----------------|----------------------------|-----------------------------|
| Group                                                                                | <i>n</i> | Pearson | Spearman | Kendall $\tau$ | Spearman CI <sub>Low</sub> | Spearman CI <sub>High</sub> |
| South                                                                                | 9        | 0.48    | 0.60     | 0.40           | -0.09                      | 0.92                        |
| Midwest                                                                              | 10       | 0.28    | 0.38     | 0.27           | -0.35                      | 0.85                        |
| West                                                                                 | 7        | 0.25    | 0.27     | 0.29           | -0.68                      | 1.00                        |
| Northeast                                                                            | 9        | 0.04    | 0.03     | -0.06          | -0.72                      | 0.63                        |
| R1                                                                                   | 29       | 0.14    | 0.10     | 0.07           | -0.26                      | 0.46                        |
| R2                                                                                   | 6        | -0.32   | -0.37    | -0.36          | -1.00                      | 0.85                        |
| Public Research                                                                      | 18       | 0.06    | -0.04    | -0.05          | -0.62                      | 0.54                        |
| Private Research                                                                     | 17       | 0.19    | 0.17     | 0.13           | -0.35                      | 0.58                        |
| Medium                                                                               | 9        | 0.31    | 0.20     | 0.14           | -0.62                      | 0.83                        |
| Large                                                                                | 21       | 0.10    | 0.04     | 0.02           | -0.46                      | 0.51                        |
| (b) Include unranked CSRankings <sub>AI</sub> by collapsing into bottom rank         |          |         |          |                |                            |                             |
| Group                                                                                | <i>n</i> | Pearson | Spearman | Kendall $\tau$ | Spearman CI <sub>Low</sub> | Spearman CI <sub>High</sub> |
| South                                                                                | 21       | 0.62    | 0.70     | 0.55           | 0.33                       | 0.87                        |
| Midwest                                                                              | 20       | 0.50    | 0.52     | 0.41           | 0.10                       | 0.79                        |
| West                                                                                 | 18       | 0.36    | 0.37     | 0.29           | -0.13                      | 0.72                        |
| Northeast                                                                            | 20       | 0.34    | 0.40     | 0.29           | -0.06                      | 0.72                        |
| R1                                                                                   | 33       | 0.22    | 0.23     | 0.15           | -0.11                      | 0.54                        |
| R2                                                                                   | 30       | -0.00   | 0.12     | 0.09           | -0.26                      | 0.48                        |
| Public Research                                                                      | 33       | 0.45    | 0.47     | 0.33           | 0.14                       | 0.70                        |
| Private Research                                                                     | 30       | 0.35    | 0.39     | 0.30           | 0.04                       | 0.67                        |
| Small                                                                                | 27       | -0.16   | -0.02    | -0.01          | -0.39                      | 0.44                        |
| Medium                                                                               | 25       | 0.45    | 0.48     | 0.37           | 0.10                       | 0.76                        |
| Large                                                                                | 27       | 0.27    | 0.25     | 0.17           | -0.17                      | 0.61                        |

Table 10. Agreement between ACAI and CSRankings<sub>AI</sub> across institutional groups with  $n > 5$ . The table reports Pearson, Spearman, and Kendall correlations. Spearman confidence intervals are based on 1000 bootstrap resamples.

| $\tau$ | Pearson $r$ | Spearman $\rho$ | Kendall $\tau$ | Spearman 95% CI | Mean $\Delta$ | $\Delta_{\text{Bottom 25\%}}$ | $\Delta_{\text{Top 25\%}}$ |
|--------|-------------|-----------------|----------------|-----------------|---------------|-------------------------------|----------------------------|
| 0.5    | 0.58***     | 0.62***         | 0.45***        | [0.45, 0.75]    | 15.65         | [1, 6]                        | [22, 72]                   |
| 1.0    | 0.55***     | 0.56***         | 0.40***        | [0.38, 0.71]    | 16.63         | [0, 6]                        | [24, 72]                   |
| 1.5    | 0.57***     | 0.54***         | 0.39***        | [0.37, 0.70]    | 17.46         | [0, 8]                        | [24, 64]                   |

Table 11. Agreement between indicator-weighted ACAI ranks and  $\text{LLM}_\tau$  ranks across sampling temperatures. Correlations are computed on the ACAI scores;  $\Delta$  is the absolute rank gap (average-tie ranking) between human and LLM. Significance: \* $p < 0.05$ , \*\* $p < 0.01$ , \*\*\* $p < 0.001$ .

## J Details on LLM-Driven Audit Study

Outputs failing schema or completeness checks were automatically retried until a valid response was produced.

Figure 12 shows the prompt used for the study.

Table 12. **Examples of Pages by Link Type (T1–T7):** This table provides illustrative examples of pages classified under each link type, chosen to help readers understand how link categories are defined and applied in the search framework.

| University                                       | Quote                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     | Link |
|--------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|
| <b>T1. University Rules/Policies/Regulations</b> |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |      |
| University of Texas at Austin                    | With the increasing integration of artificial intelligence (AI) tools—such as ChatGPT, Copilot, Bard Gemini, Claude and other generative AI applications known as large language models (LLM), diffusion models, or generative AI applications—into university activities, it's essential to use these technologies responsibly. This guidance, developed collaboratively by the Office of Legal Affairs, University Compliance Services, the Information Security Office, and the Business Contracts Office, outlines acceptable practices for utilizing generative AI tools while safeguarding institutional, personal, and proprietary information. Additional guidance may be forthcoming as circumstances evolve.                    | Link |
| Georgia Southern University                      | This policy establishes guidance for the responsible, ethical, and transparent use of Artificial Intelligence (AI) tools at Georgia Southern University (University) including in teaching, learning, assessment, classroom activities, University community service, research, creative activity, scholarly communication, and administrative activities, while encouraging innovation, academic freedom, and appropriate autonomy while maintaining compliance with all data security and privacy regulations. This policy ensures that the University complies with Board of Regents (BOR) Policy 6.28 Artificial Intelligence in Academic Context; federal, state, and international laws; and industry standards and best practices. | Link |
| Carleton College                                 | AI technologies raise novel questions around data security, attribution, and ethics. In many cases, Carleton's existing policies still apply to AI, but in some cases this technology requires new policies or new interpretations of existing policies. This page will provide links to policies that apply directly to AI use at Carleton and highlight any additions and changes as they're made.                                                                                                                                                                                                                                                                                                                                      | Link |
| University of Southern California                | This Research Guide provides information on the use of Generative AI in academic papers and research, and provides guidance on the ethical use of Generative AI in an academic setting.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Link |
| <b>T2. Center for Teaching &amp; Learning</b>    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |      |
| Stanford University                              | AI Meets Education at Stanford (AIMES) is a VPUE effort to catalyze and support critical engagement with generative AI in Stanford teaching and learning contexts, coordinated by the Center for Teaching and Learning.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Link |

*Continued on next page*

| University                                         | Quote                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  | Link                 |
|----------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|
| Ball State University                              | Explore a variety of courses designed to meet learners at all levels. Whether you're looking for introductory classes or advanced specialization tracks, these courses provide structured, in-depth instruction in AI topics to build your skills and confidence. A strong grasp of AI terminology is essential for navigating complex concepts and discussions. This glossary of terms offers definitions and explanations of key terms, serving as a quick reference to clarify AI language as you advance in your learning.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | <a href="#">Link</a> |
| University of Michigan at Ann Arbor                | The release of ChatGPT in late 2022 jump started an ongoing and growing exchange in higher education about both the promises and significant risks posed by Generative Artificial Intelligence, particularly to the teaching and learning enterprise. This site is designed to offer links to programs and resources from U-M and beyond to help you navigate this new landscape. Given how rapidly the GenAI landscape is shifting, we include links to sources that offer regular posts and updates on this topic.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | <a href="#">Link</a> |
| University of Florida                              | How is Artificial Intelligence (AI) affecting teaching and learning in higher education? Artificial intelligence (AI) is significantly impacting higher education, revolutionizing various aspects of the learning experience. AI-powered tools and platforms are transforming how students access educational content, tailor their learning paths, and receive personalized feedback. Moreover, AI-driven systems can help educators generate educational content and facilitate research endeavors. While these advancements bring exciting opportunities, it's essential to address ethical concerns, data privacy, and ensure AI complements the education process rather than replacing it entirely. Embracing AI responsibly can lead to a more accessible, efficient, and effective higher education landscape. Artificial intelligence has introduced significant challenges to academic integrity in education. As AI becomes more accessible, educators have expressed concerns about students using it to generate answers to questions on tests and assignments. Rather than reacting in fearful ways to new advances in AI, educators can focus on potential benefits, such as providing new perspectives on a problem and generating content that can be analyzed or critiqued. Undoubtedly, faculty need to provide guidelines to students about the appropriate and inappropriate uses of AI tools. However, faculty can also model and encourage productive and positive uses of AI and help students see its value. | <a href="#">Link</a> |
| <b>T3. AI Institute/Initiative/Center/Hub/Al@U</b> |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |                      |
| University of Wyoming                              | The University of Wyoming's AI Initiative is a bold, people-centered effort to shape the future of our state, empowering citizens and communities to thrive in an AI-driven world. By addressing key industries like agriculture, engineering, energy, tourism, wildlife conservation, and rural healthcare, UW is ensuring that AI enriches lives and drives sustainable growth. This initiative will enhance the University of Wyoming's ability to bring advances in AI to disciplines across the university to advance the state. It will attract investments, build corporate partnerships, seed entrepreneurship, and equip every student and community to participate in the global AI transformation, securing a prosperous future for all of Wyoming.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | <a href="#">Link</a> |
| University of California, Berkeley                 | Welcome to the AI Hub – your central resource for artificial intelligence at UC Berkeley. Rooted in Berkeley's pioneering spirit, ethos of inclusivity, and culture of excellence, this hub connects our community with essential AI tools, training, policies, and opportunities. Whether you're a student, researcher, faculty, or staff member, you'll find guidance, collaboration, and innovation here to help navigate the evolving world of AI and amplify the impact of our collective efforts.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | <a href="#">Link</a> |

*Continued on next page*

| University                                  | Quote                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | Link                 |
|---------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|
| San José State University                   | Welcome to the bold new world of Artificial Intelligence (AI), where groundbreaking innovation meets inclusive leadership in the heart of Silicon Valley. At San José State University, we are a place of firsts, pioneering advancements in AI and empowering the next generation of leaders to shape the future of technology. Our interdisciplinary programs blend cutting-edge research with hands-on learning, equipping students and professionals to solve real-world challenges and explore ethical solutions in AI. As a proud partner in Silicon Valley's ecosystem of global innovation, SJSU connects you to industry leaders, transformational opportunities, and a vibrant community. Whether you're forging new paths in AI development, exploring its societal impact or preparing to lead in this dynamic field we'll help you unlock your potential and create a future where everyone can thrive. San José State University leads the way in AI innovation and leadership. | <a href="#">Link</a> |
| Clark Atlanta University                    | The NSF Expand-AI project led by Clark Atlanta University (CAU) in collaboration with AI4OPT builds an AI Hub at CAU to transform accessibility to AI jobs, AI research, and the AI ecosystem. AIHub@CAU consists of three key pillars: (1) a Master Program in AI; (2) a PhD program in AI; and (3) research collaborations between CAU and Georgia Tech. The program is a joint project by the department of mathematical sciences, the department of Cyber-Physical Systems and the School of Business Administration, making it truly multidisciplinary.                                                                                                                                                                                                                                                                                                                                                                                                                                  | <a href="#">Link</a> |
| <b>T4. Library/Library Guides</b>           |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |                      |
| Mercer University                           | Definitions. Artificial Intelligence: AI is typically defined as the ability of a machine to perform cognitive functions we associate with human minds, such as perceiving, reasoning, learning, and problem solving. Examples of technologies that enable AI to solve complex problems include robotics, computer vision, language, virtual agents and machine learning.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     | <a href="#">Link</a> |
| Chapman University                          | AI Literacy. Use this guide to understand Artificial Intelligence literacy in the context of higher education. What is AI Literacy? "AI literacy is the ability to understand, use, and think critically about AI technologies and their impact on society, ethics, and everyday life." - Lo, L. S. (2025). AI Literacy: A Guide for Academic Libraries. College & Research Libraries News, 86(3), Article 3. <a href="https://doi.org/10.5860/crln.86.3.120">https://doi.org/10.5860/crln.86.3.120</a>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | <a href="#">Link</a> |
| University of Chicago                       | Generative AI. Information about Generative AI tools and their use in and outside of the classroom.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           | <a href="#">Link</a> |
| Brown University                            | Generative artificial intelligence has already started to have an impact on the way we discover, manage, create, and disseminate information. Generative AI tools are in a state of rapid development, and new information about applications, policies, and social impact is released each day. While every attempt will be made to keep this guide up to date, please be aware that the information included here is likely to age quickly. This guide includes context and advice for engaging with generative Artificial Intelligence, and does not represent University policy. The Library does not endorse any specific AI technologies, and encourages users to be cautious about sharing personal information when using AI tools.                                                                                                                                                                                                                                                   | <a href="#">Link</a> |
| <b>T5. Academic Integrity/Honor Code</b>    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |                      |
| Creighton University                        | Cheating: The deliberate use or attempted use of unauthorized material in an academic exercise, including unauthorized collaboration with classmates, or use of unauthorized work created by artificial intelligence.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | <a href="#">Link</a> |
| Wofford College                             | Unauthorized use of generative artificial intelligence to create content that is submitted as one's own.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      | <a href="#">Link</a> |
| Marquette University                        | Academic Integrity. A Message for Faculty, Staff, and Students on the use of Large Language Model-Based Chatbots ("generative artificial intelligence") at Marquette University.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              | <a href="#">Link</a> |
| University of New Hampshire                 | Cheating. Use or attempted use of any academic exercise materials, information, study aids, electronic data, AI tools, assignment/exam surrogate, or other forms of assistance without authorization.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | <a href="#">Link</a> |
| <b>T6. AI Steering Committee/Task Force</b> |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |                      |

*Continued on next page*

| University                              | Quote                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     | Link                 |
|-----------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|
| Texas A&M University                    | The Artificial Intelligence, Innovative & Emerging Technologies Work Group equips faculty with the knowledge and resources to explore and integrate cutting-edge technologies into teaching and learning. By curating and sharing AI-related resources, fostering collaboration, and promoting best practices, the group empowers educators to leverage emerging technologies to enhance student engagement and academic success across the A&M System.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | <a href="#">Link</a> |
| Stony Brook University                  | Library AI Steering Committee. The Stony Brook University Libraries AI Steering Committee plays a central role in guiding the responsible and strategic integration of artificial intelligence across library services, operations, and research support. Established to ensure that emerging technologies advance—rather than compromise—the Libraries’ core values of equity, accessibility, intellectual freedom, and responsible innovation, the committee evaluates opportunities and risks, recommends best practices, and supports evidence-based decision-making. Its charge includes reviewing AI initiatives for alignment with institutional priorities, developing ethical guidelines and principles, promoting staff training and AI literacy, and fostering collaborations with campus partners and professional communities. Through regular reporting and transparent communication with library leadership, the AI Steering Committee helps ensure accountability and positions the Libraries to thoughtfully and proactively navigate the evolving landscape of AI in higher education. | <a href="#">Link</a> |
| California State University, Long Beach | The purpose of the AI Academic Subcommittee is to explore AI technologies and plan for future implementations. The subcommittee will make recommendations to the AI Steering Committee to develop guidelines for campus-wide deployment. To foster a community of AI users on campus, the subcommittee will also make recommendations for professional development and support for faculty and staff on AI-related topics.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | <a href="#">Link</a> |
| Iowa State University                   | The 2024 Generative AI Guidance Committee has successfully completed its charge. The subcommittees were assembled, carried out their tasks diligently, and contributed valuable insights and recommendations. Their work has laid a strong foundation for our institutional AI strategy. For more information and access to resources, please visit <a href="https://ai.iastate.edu">ai.iastate.edu</a> .                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | <a href="#">Link</a> |
| <b>T7. Other Relevant Links</b>         |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |                      |
| University of Michigan at Ann Arbor     | Custom GenAI Services for the U-M Community. U-M is proud to be the first university in the world to provide a custom suite of generative AI tools to its community. With a focus on equity, accessibility, and privacy, our AI Services are available to all U-M faculty, staff, and students on the Ann Arbor, Flint, Dearborn, and Michigan Medicine campuses.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | <a href="#">Link</a> |
| University of Chicago                   | Advances in artificial intelligence (AI) and data science are driving breakthroughs – transforming scientific discovery, accelerating innovation, and changing entire industries. At the University of Chicago, our long-standing tradition of rigorous inquiry and interdisciplinary collaboration provides a powerful foundation for tackling large-scale problems in AI and data science and unleashing their greatest potential to improve individual lives and our world. Faculty and researchers from every division and school at UChicago are at the forefront of these fields, from developing trustworthy AI systems and foundational frameworks to enabling transformative advances in areas such as precision medicine and next-generation climate modeling.                                                                                                                                                                                                                                                                                                                                  | <a href="#">Link</a> |
| Northern Illinois University            | NIU has introduced Mission, a new feature for undergraduate students, that includes reaching out to you through AI-assisted text messages and a chatbot to respond to your questions around the clock. Mission will help answer your questions, and connect you to campus services and information, including: Academic success - tutors, academic advisors, study skills and more. Financial matters - financial aid, FAFSA and more. Student life and involvement - clubs, organizations, events and more. Well-being/mental and physical health - counseling services, nutrition, campus recreation and more. Mission will provide timely, accurate responses via text to your questions at all times of day, regardless of your location. No logins or app downloads are required.                                                                                                                                                                                                                                                                                                                    | <a href="#">Link</a> |

*Continued on next page*

| University                      | Quote                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      | Link                 |
|---------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|
| Case Western Reserve University | University Technology offers many services and applications related to Generative AI. Below are some AI technologies available to the campus community. Note: Consumer AI services, especially free ones, often collect the data you enter into them and use that data in their training models. This can lead to your data being made available via these AI services. Never put sensitive university information into an AI service if the university does not have a contract with the AI vendor with proper privacy and security safeguards. The university offers AI services that will protect your data. Use those services when sensitive information is involved. | <a href="#">Link</a> |

**Privacy Policy****Last Updated:** December 29, 2025**Platform:** Prolific (<https://www.prolific.com>)If you have questions about this study, you may contact the research team at [redacted@redacted](mailto:redacted@redacted).**1. Purpose of the Task**

You are being asked to participate as an **annotator** in a research project examining how U.S. universities define and respond to **artificial intelligence (AI)** in their institutional policies and teaching resources. Your role involves reviewing university web pages and coding the presence or absence of certain indicators related to AI governance, academic integrity, and instructional support. The study analyzes publicly available institutional documents and does not evaluate individual instructors, students, or staff.

This task contributes to a larger academic study focused on understanding patterns in higher education responses to AI technologies.

**2. Data We Collect**

During your participation, we collect the following categories of data:

- **Prolific ID and Prolific-provided Demographic Data:** This data is provided by the Prolific platform to us. Your Prolific ID will be used only for compensation and quality control, and will be removed as part of the anonymization process before dataset release.
- **Timing Metadata:** The start and completion time for your task submission, which helps assess annotation duration and data quality.
- **Annotation Data:** Your coded responses and URLs that you provide.

No personal browsing history, IP address, or system-level data is collected by the researchers; such information remains with Prolific and is governed by their Privacy Policy.

**3. How Your Data Is Used**

Your coded responses will be:

- **De-identified.**
- **Used for academic research and publication** in peer-reviewed journals or conference presentations.

**De-identified datasets will be shared publicly for use by other researchers under ethical data-sharing agreements, consistent with open science practices.**

**4. Data Storage and Security**

All data collected will be stored securely in **encrypted storage** (e.g., Google cloud) accessible only to the research team.

- Data will be **retained for up to 3 years** after study completion and then deleted or permanently anonymized.
- No data will be sold or shared with commercial entities.

**5. Voluntary Participation and Withdrawal**

Your participation is **entirely voluntary**. You may withdraw from the task at any time prior to submission on Prolific. If you withdraw before completing the task, no partial data will be used.

**6. Risks and Benefits**

There are **minimal risks** associated with participation. You may experience minor fatigue from reviewing university materials. There are **no direct personal benefits**, though your work contributes to research improving understanding of AI use in education policy.

**7. Confidentiality**

Your responses will never be linked to your name or contact information. Any publications or presentations resulting from this research will contain only **anonymous findings**.

**8. Consent**

By completing the task on Prolific, you confirm that you:

- Are 18 years of age or older.
- Understand the nature and purpose of this research.
- Consent to your **anonymized** responses being used for research and publication purposes.

Fig. 9. Privacy Policy Provided to Human Annotators Recruited via the Prolific Platform.




Fig. 10. **Spearman  $\rho$  rank correlations between ACAI ranks and CSRankings<sub>AI</sub> ranks across institutional subgroups show that AI governance capacity and AI research output are almost entirely uncorrelated ( $\triangleright F3$ ):** All CIs cross zero, unranked CSRankings<sub>AI</sub> universities are excluded, and categories with  $n \leq 5$  are dropped. We provide detailed results in Table 10.

## CSRankings: Computer Science Rankings




Fig. 11. **CSRankings Website:** Used to approximate research output, accessible via <https://csrankings.org>. Retrieved 1/1/26.

Evaluating whether a university meets specific criteria based strictly on the provided institutional links.

University: {UNIVERSITY}

Approved sources (you must not use any external sites, but you should explore sublinks from these sites):

{LINK\_LIST}

Evaluation criteria:

{CRITERIA\_LIST}

Scoring scale (select exactly one per criterion):

A. Present/Yes – A clear statement directly addressing the item is found on an institutional page within 5 minutes.

B. Partial/Implicit/Somewhat – The item is mentioned or implied, but key details are missing.

C. Absent/No – You reasonably searched the allowed sources and did not find relevant content.

D. Unclear or Took Longer Than 5 Minutes – Navigation difficulty, vague language, or time limits prevented a confident decision.

E. Conflicting Information – Different institutional sources provide contradictory guidance for the same item.

Output requirements:

- Return VALID JSON ONLY.
- No markdown, no commentary.
- Return an ARRAY with one object per criterion.
- Every criterion MUST appear exactly once.

Schema:

```
[
  {
    "criterion": "A1",
    "score": "A|B|C|D|E",
    "urls": ["https://...", "https://..."]
  }
]
```

Fig. 12. **Prompt used in the LLM robustness study.** The figure shows the full evaluation prompt provided to the language model, including source restrictions, scoring criteria, and the required JSON output schema, mirroring the instructions given to human annotators in the ACAI audit.